



A stream processing pipeline for an online advertising platform

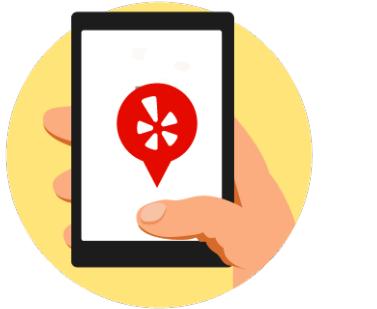
Yifan Wang

Amit Ramesh



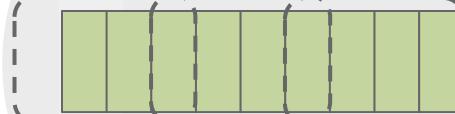
Yelp's Mission

Connecting people with great
local businesses.



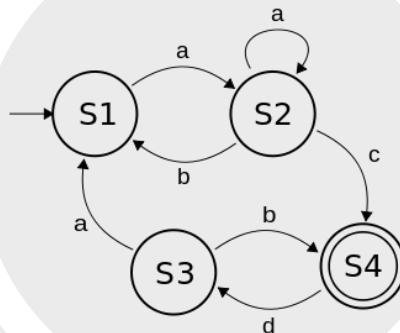


Yelp Ads



Sliding Window

State Management



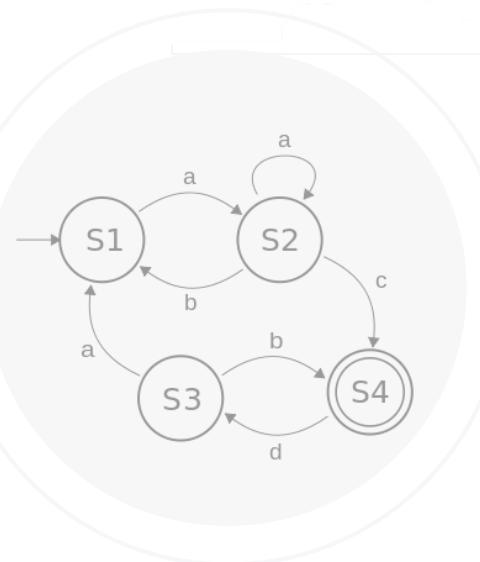


Yelp Ads

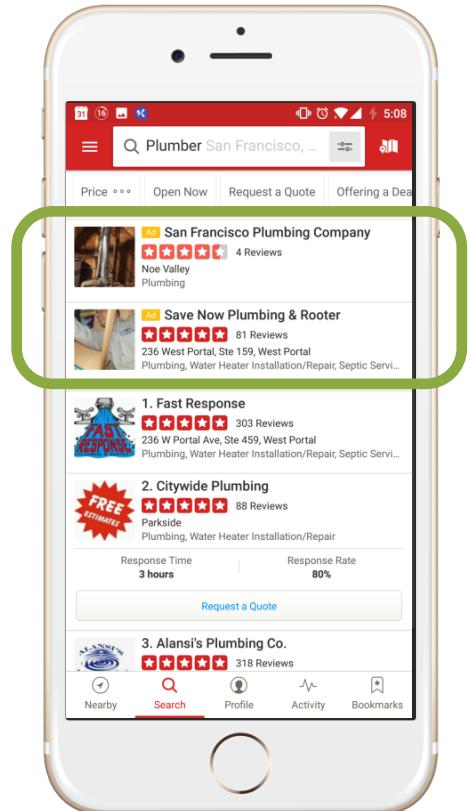
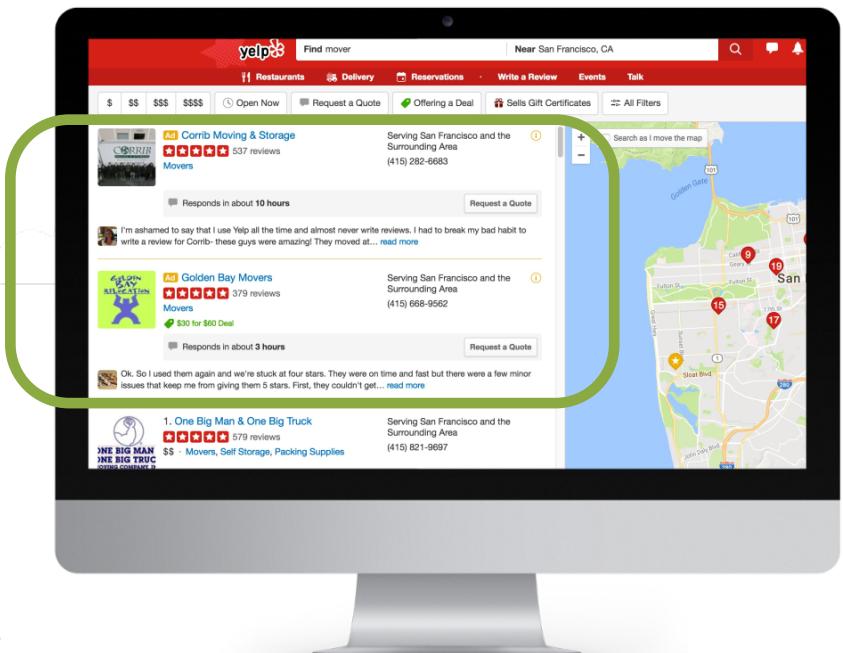


Sliding Window

State Management



Yelp Ads



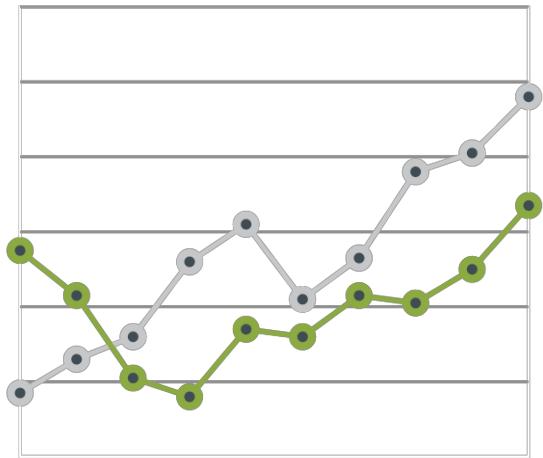
Events Timeline



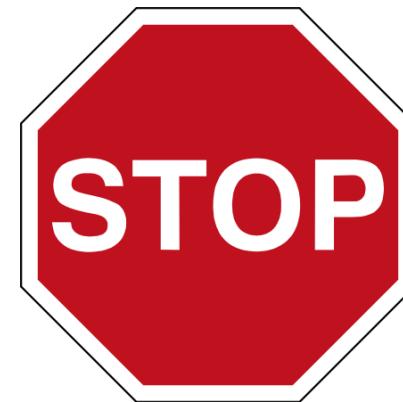
- Opportunity
- View
- Click



Why Real-time?



Real-time Metrics



Campaign Overspending



Overall Process

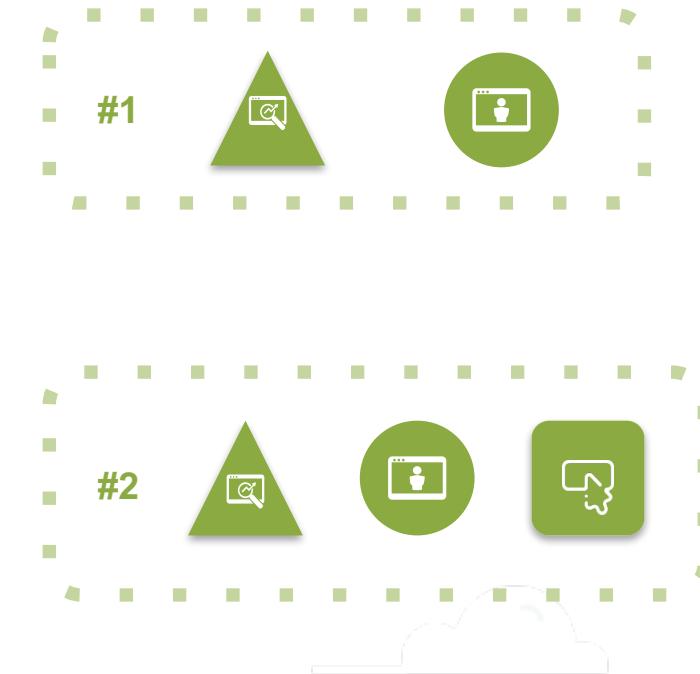
Input from Kafka



Overall Process



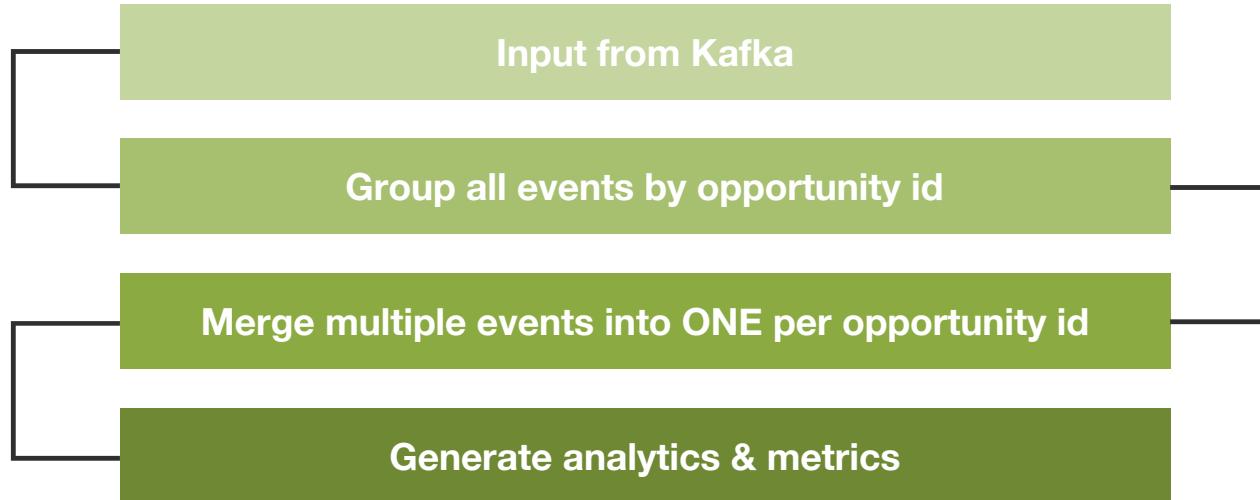
Group Events By Id



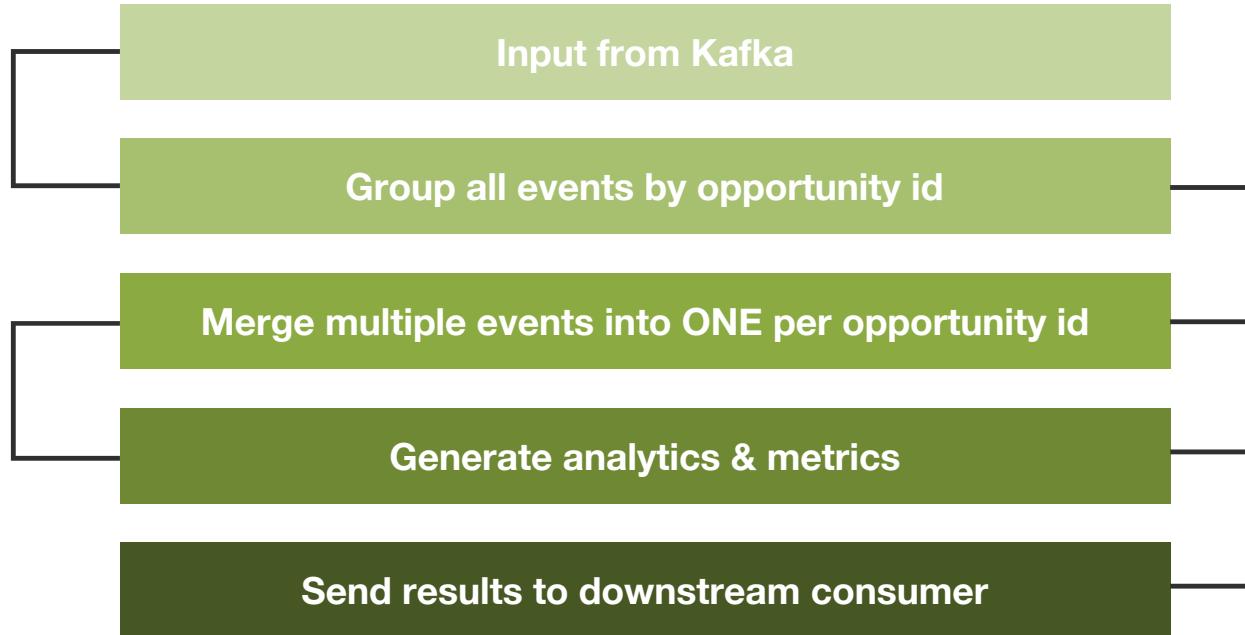
Overall Process



Overall Process

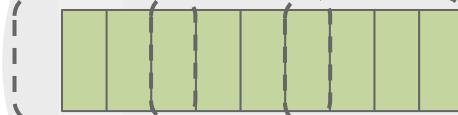


Overall Process



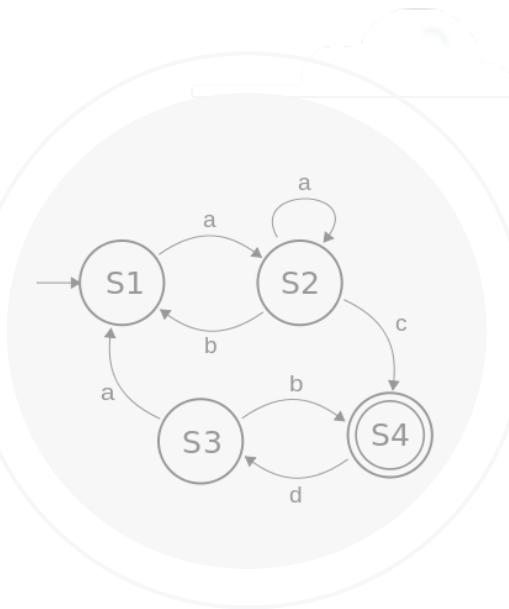


Yelp Ads

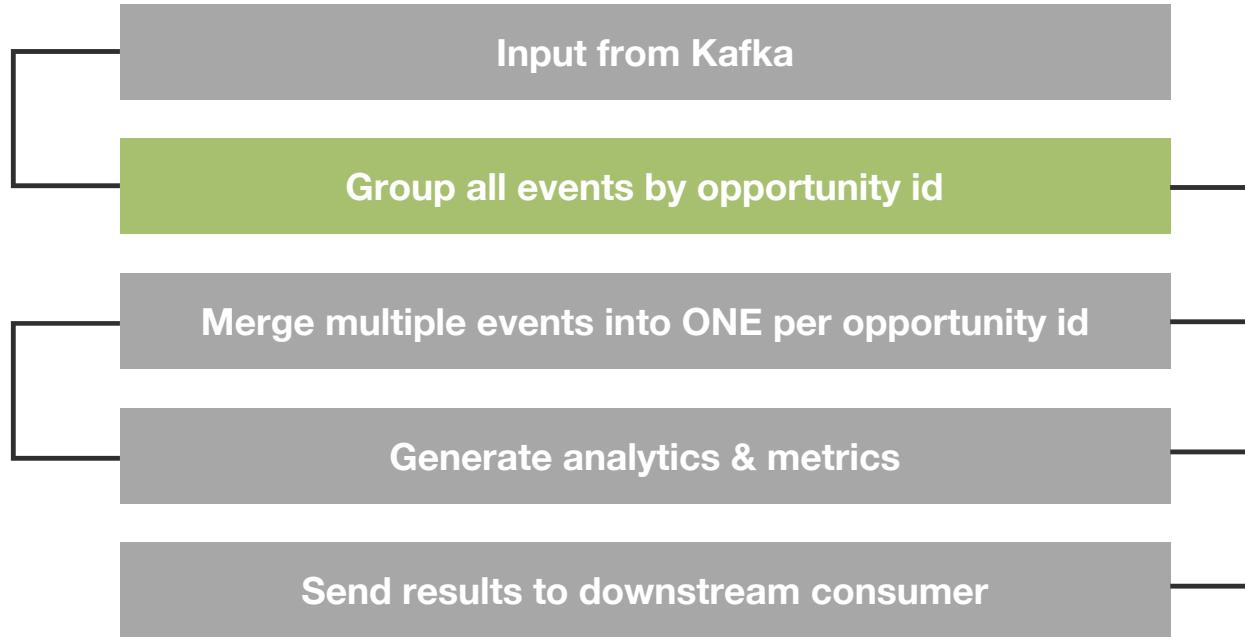


Sliding
Window

State
Management



Overall Process



Sliding Window in Spark



Spark API:

- `Window()`
- `countByWindow()`
- `reduceByWindow()`
- `reduceByKeyAndWindow()`
- `countByValueAndWindow()`



Sliding Window in Spark



?

Challenges with existing API



No state tracking



Do not support complex customized business logic



Sliding Window



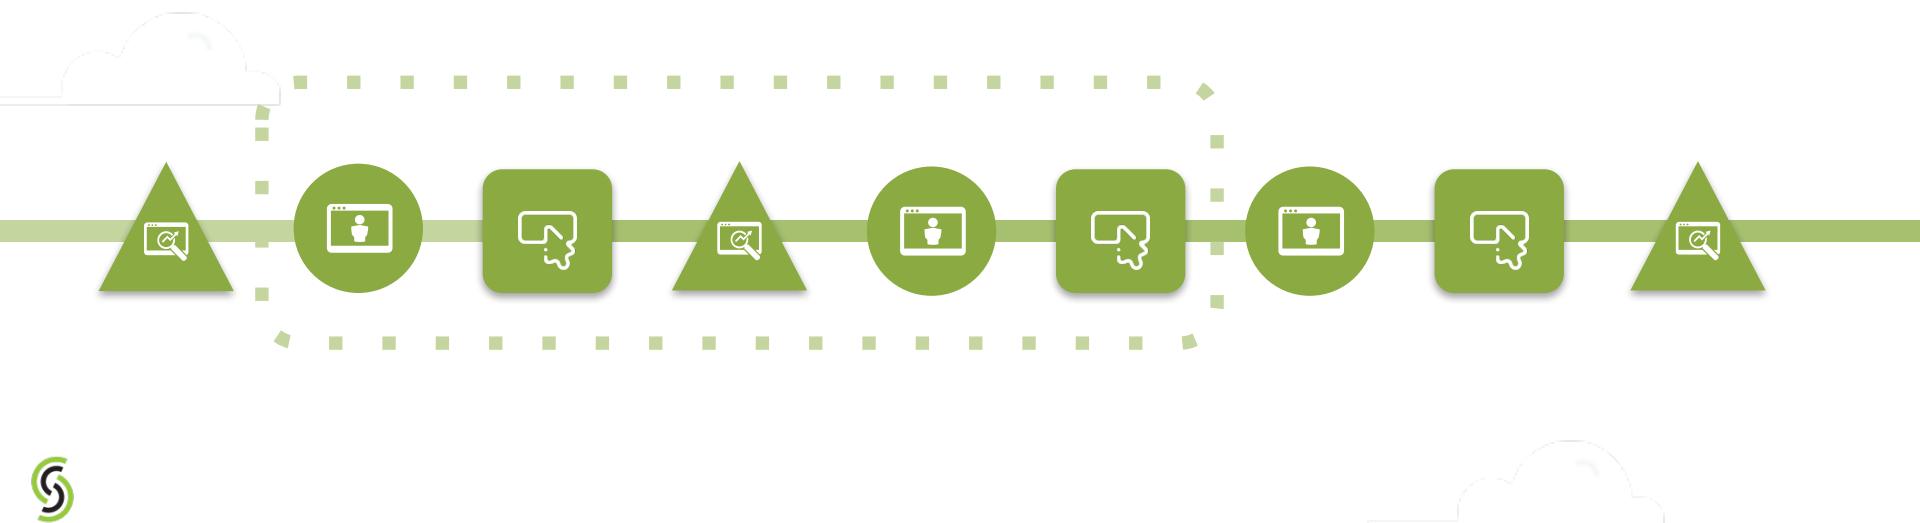
Opportunity



View



Click



Sliding Window



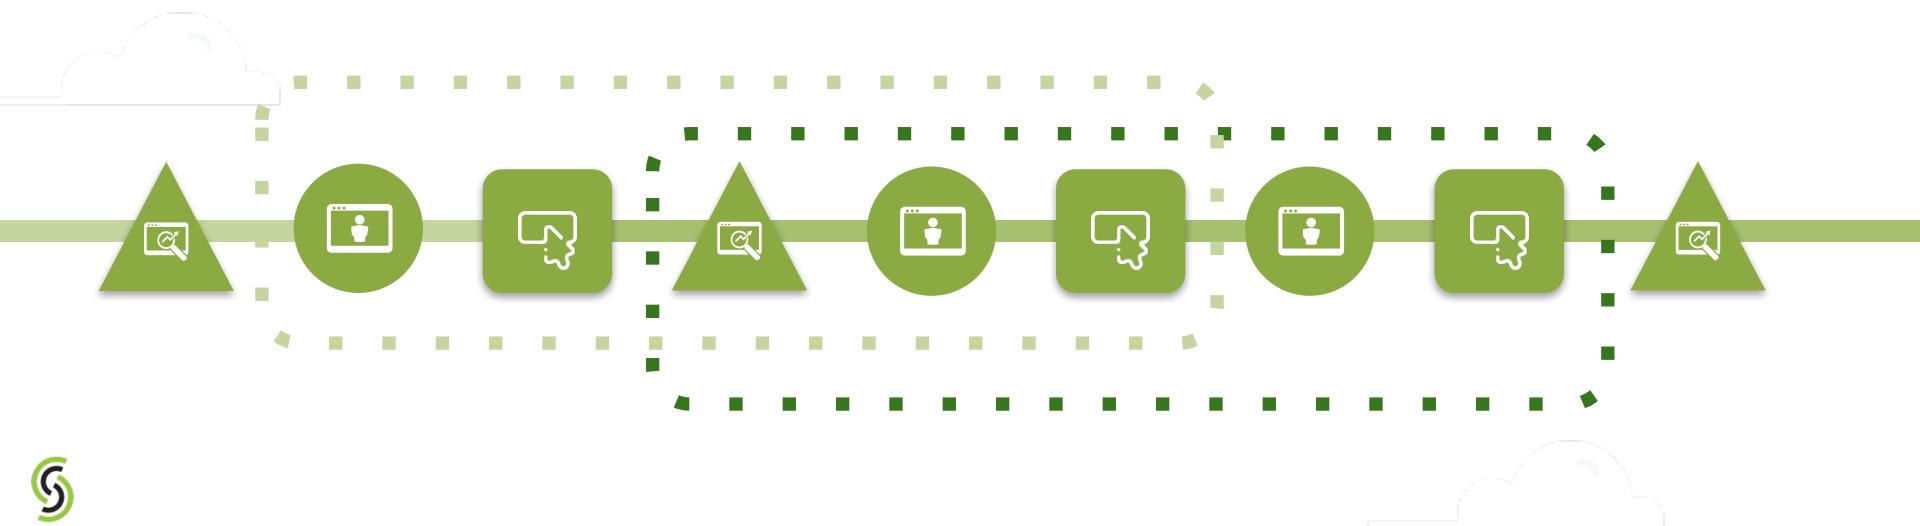
Opportunity



View



Click



Sliding Window in Spark

?

Challenges with existing API

- No state tracking
- Do not support complex customized business logic

Sliding Window



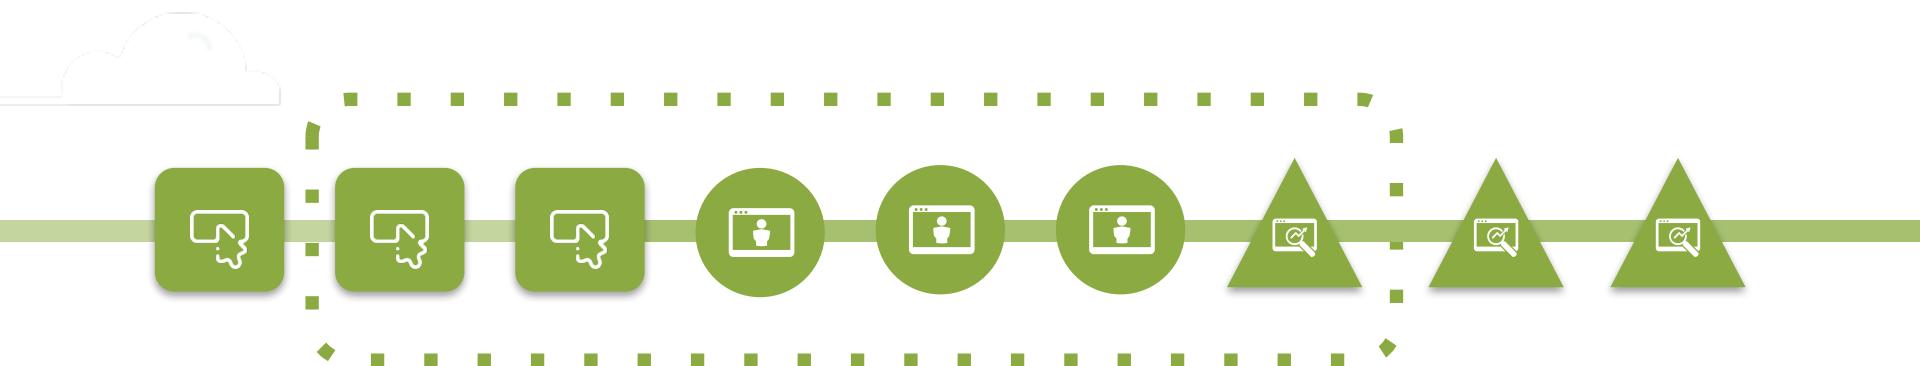
Opportunity



View



Click



Our solution

UpdateStateByKey(update_function)



Our solution

`UpdateStateByKey(update_function)`



Ability to apply complex
business rules to ad events



Our solution

`UpdateStateByKey(update_function)`

Ability to apply complex
business rules to ad events

Build-in mechanism
to track states for
any key-value pair



Our solution



1

2

3



Attach expire date/time
when events are first seen
& state is initialized

Drop the state if it expires

Apply business logic to
new events + current state

Our solution



1

2

3



Attach expire date/time
when events are first seen
& state is initialized

Drop the state if it expires

Apply business logic to
new events + current state

Our solution



1

2

3



Attach expire date/time
when events are first seen
& state is initialized

Drop the state if it expires

Apply business logic to
new events + current state

Pseudo code:

```
def update_function(new_events, current_state):  
    if current_state is None:  
        current_state = init_state()  
        attach_expire_datetime(current_state)
```



1

Pseudo code:

```
def update_function(new_events, current_state):  
    if current_state is None:  
        current_state = init_state()  
        attach_expire_datetime(current_state)  
  
    .....  
    if is_expired(current_state):  
        return None
```

1

2

Pseudo code:

```
def update_function(new_events, current_state):  
  
    if current_state is None:  
        current_state = init_state()  
        attach_expire_datetime(new_events)  
  
    .....  
  
    if is_expired(current_state):  
        return None  
  
    .....  
  
    if new_events:  
        apply_business_logic(new_events, current_state)
```

1

2

3

Downside

Spark will apply the state update function for all existing keys

Return early in your `update_function()`



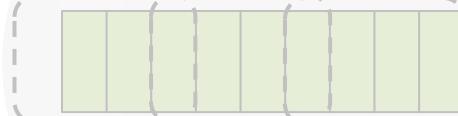
Improvement in Spark 1.6+



<https://databricks.com/blog/2016/02/01/faster-stateful-stream-processing-in-apache-spark-streaming.html>

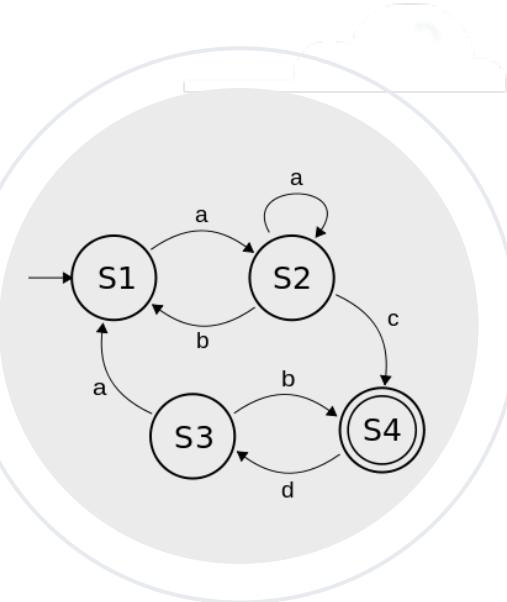


Yelp Ads

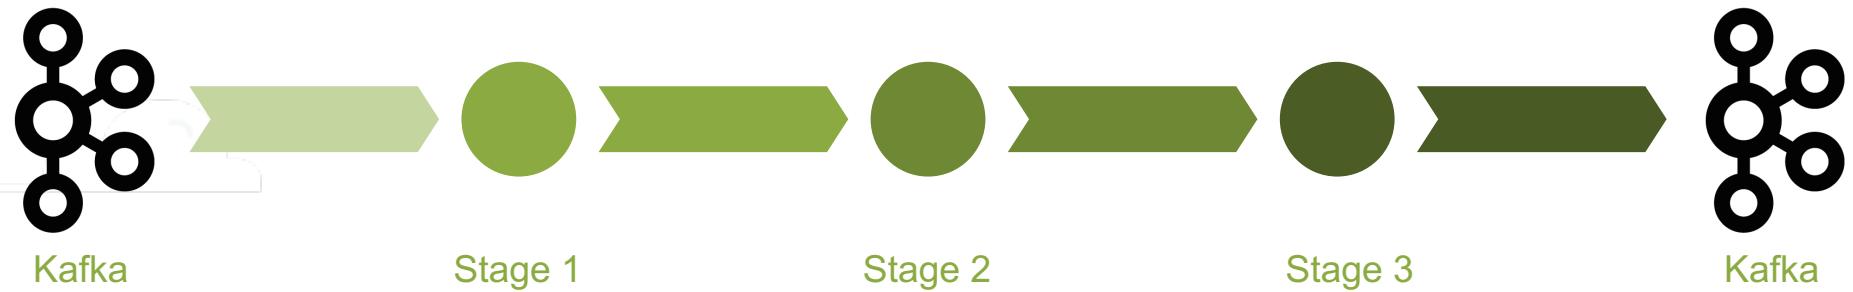


**Sliding
Window**

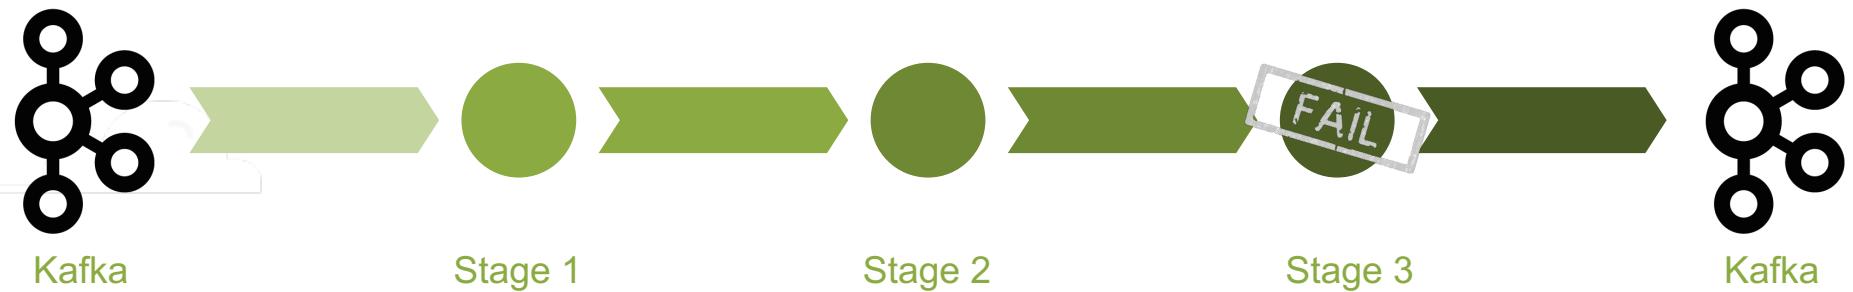
**State
Management**



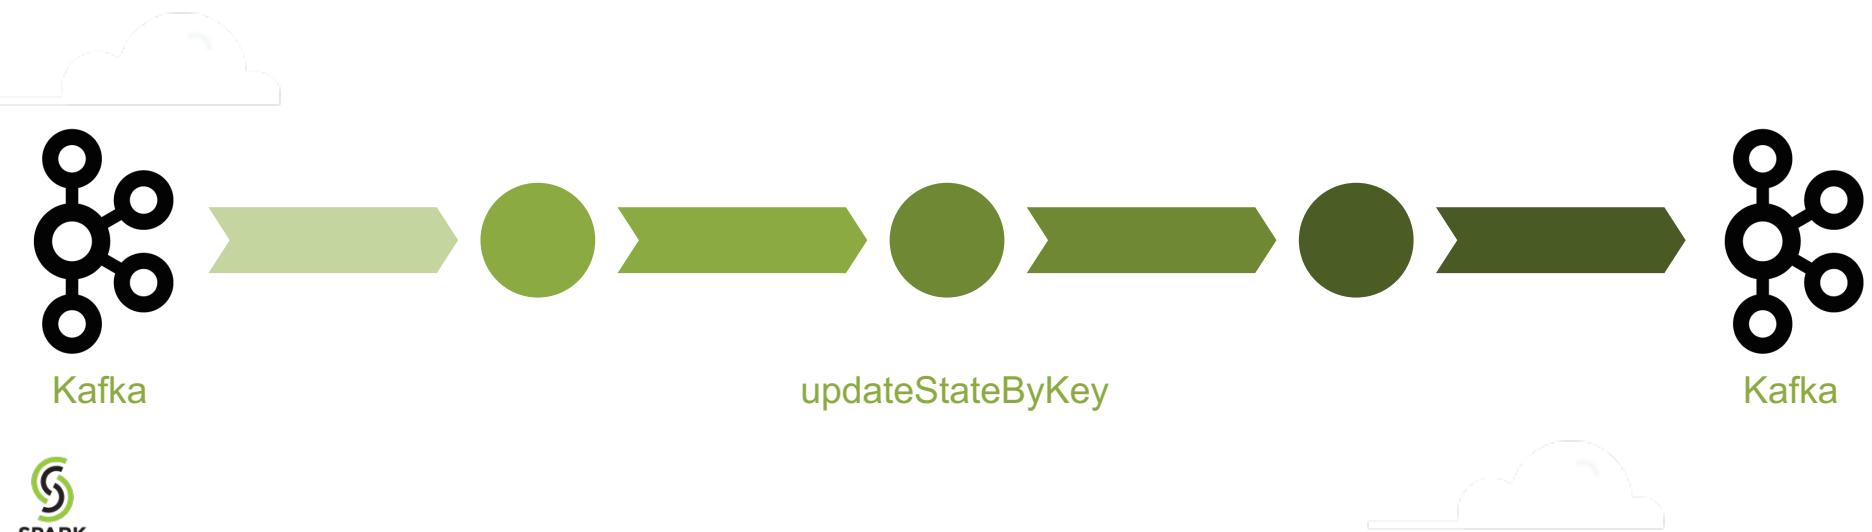
Dependency graph



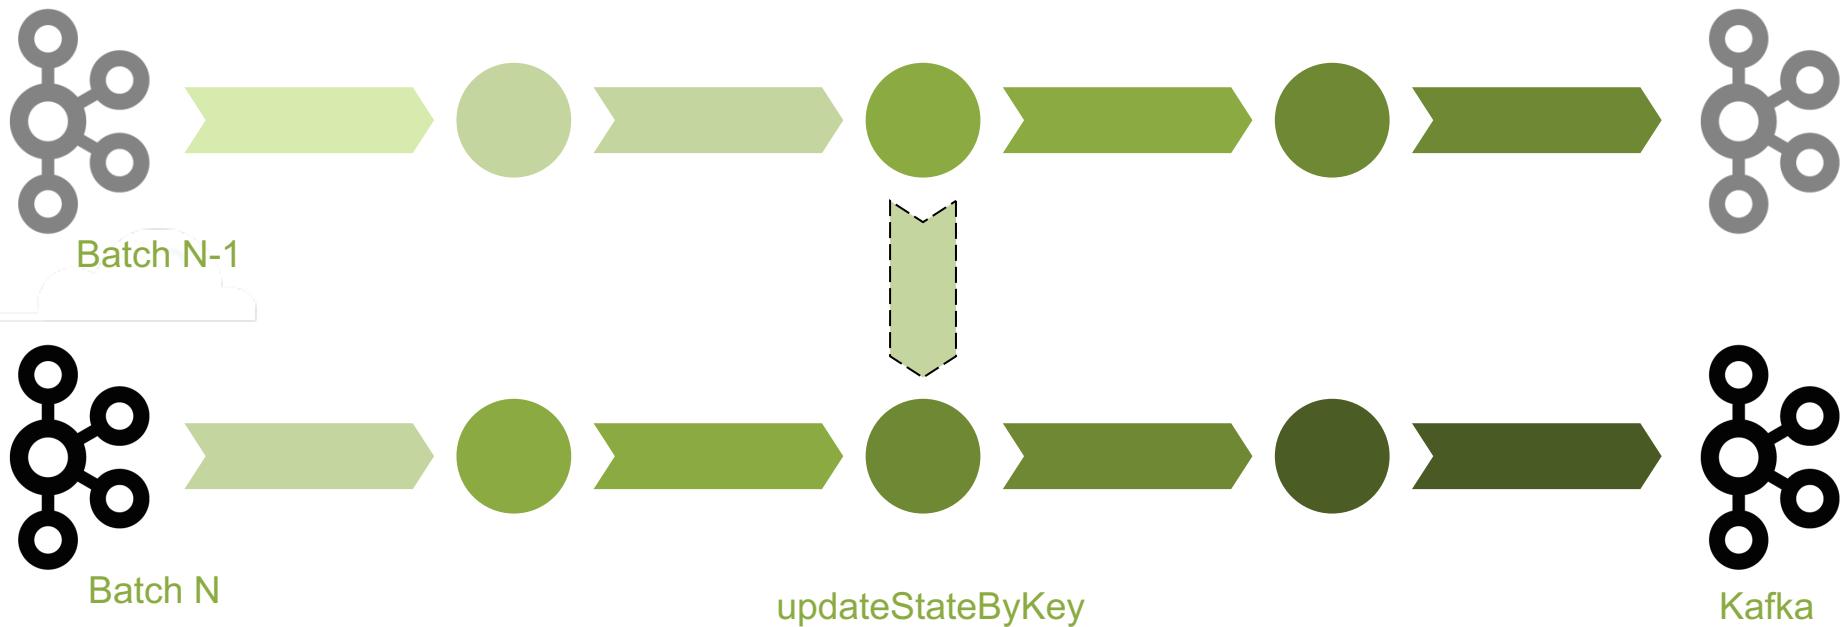
Dependency graph



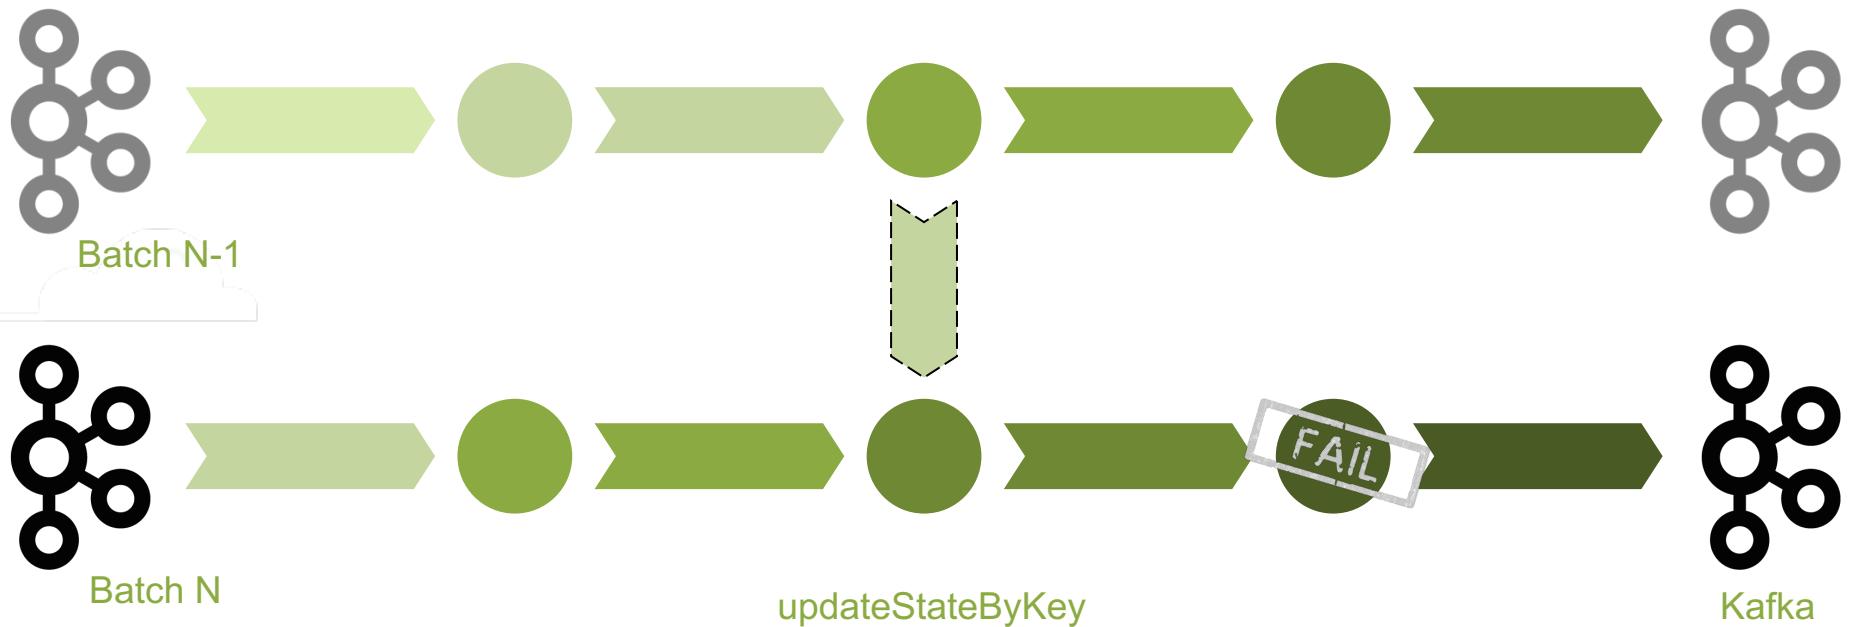
Dependency graph (stateful)



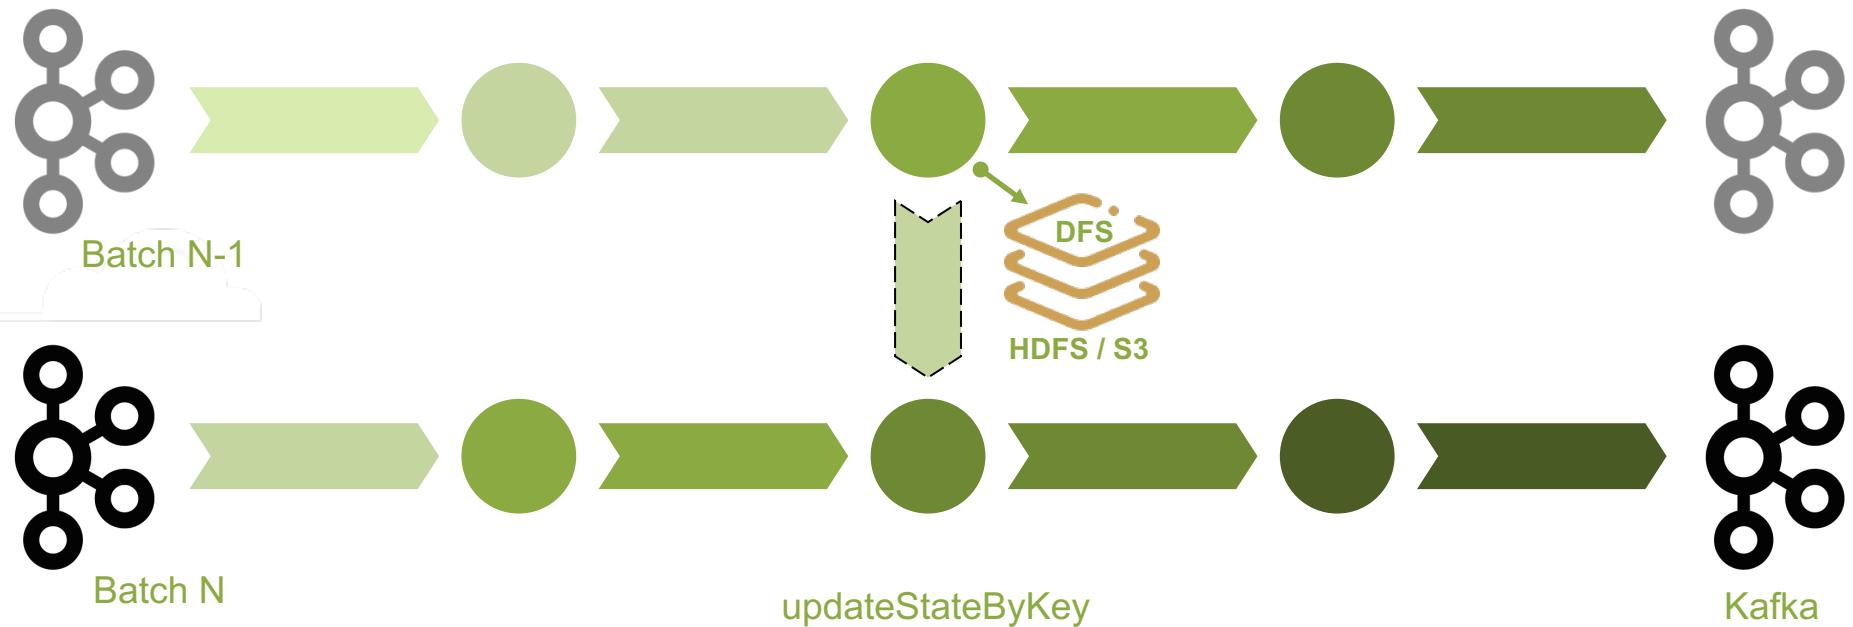
Dependency graph (stateful)



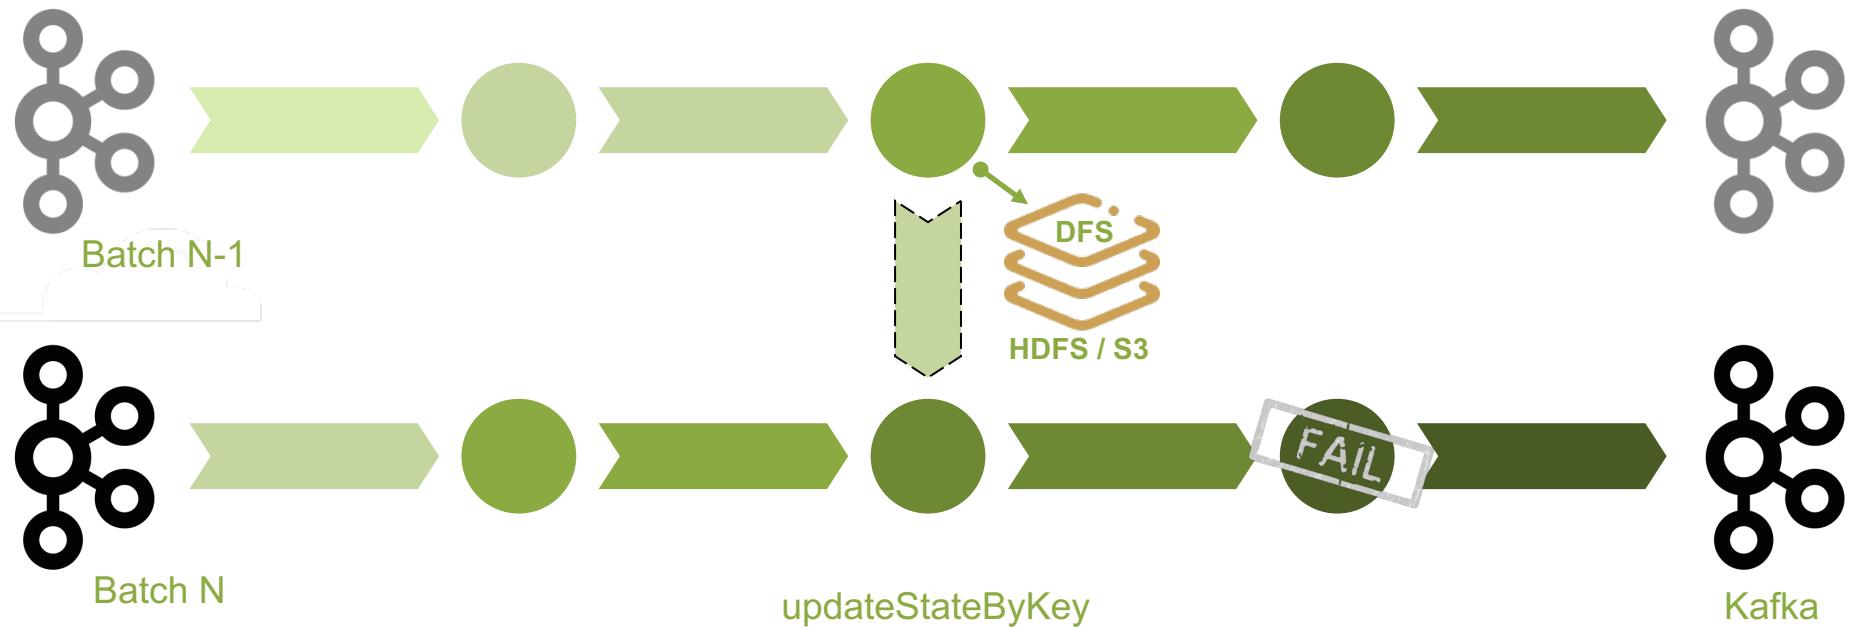
Dependency graph (stateful)



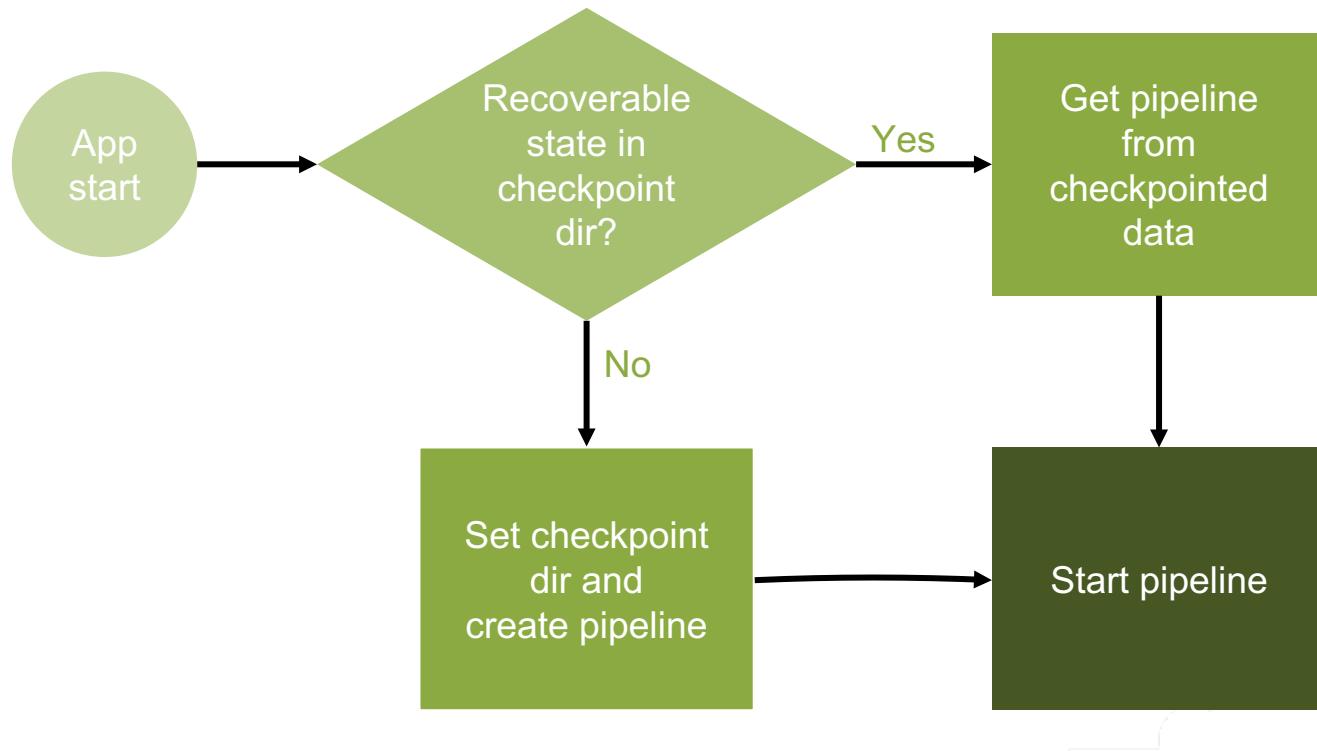
Checkpointing



Checkpointing



Checkpoint recovery

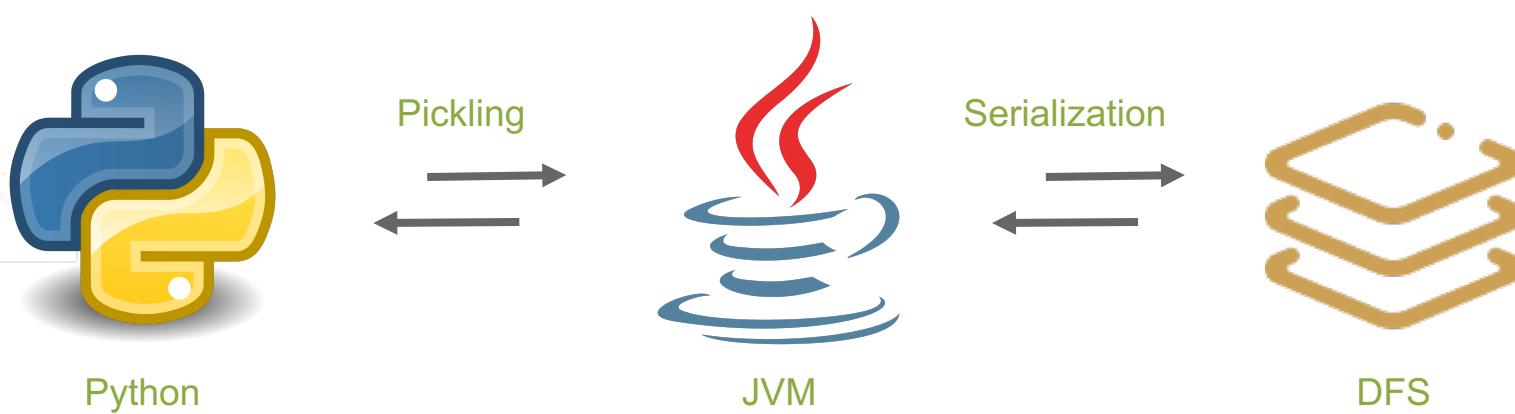


```
if __name__ == '__main__':
    ssc = StreamingContext.getOrCreate(CHECKPOINT_DIR, create_ssc)
    ssc.start()
    ssc.awaitTermination()
```

```
if __name__ == '__main__':
    ssc = StreamingContext.getOrCreate(CHECKPOINT_DIR, create_ssc)
    ssc.start()
    ssc.awaitTermination()

def create_ssc():
    sc = SparkContext()
    ssc = StreamingContext(sc, BATCH_INTERVAL)
    ssc.checkpoint(CHECKPOINT_DIR)
    setup_pipeline(ssc)
    return ssc
```

Checkpointing and recovery



Restoring state via data replay



Replay how much?

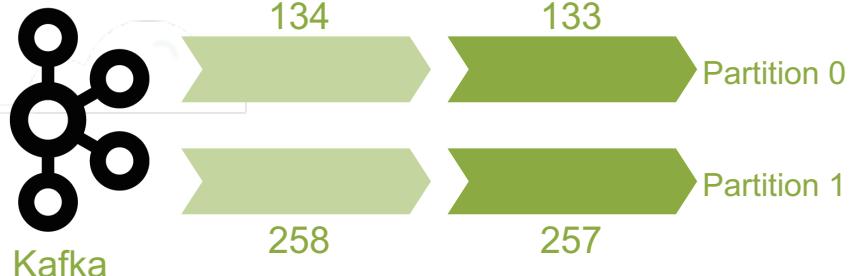


updateStateByKey

```
{  
  'key_1': [  
    {'event': 'event_15'},  
    {'event': 'event_20'},  
  ],  
  'key_2': [  
    {'event': 'event_23'},  
    {'event': 'event_13'},  
  ],  
}
```



Replay from where?



updateStateByKey

```
{  
  'key_1': [  
    {'event': '...', 'partition': 0, 'offset': 130},  
    {'event': '...', 'partition': 1, 'offset': 230},  
  ],  
  'key_2': [  
    {'event': '...', 'partition': 0, 'offset': 119},  
    {'event': '...', 'partition': 1, 'offset': 251},  
  ],  
}
```



Externalizing partial state



updateStateByKey

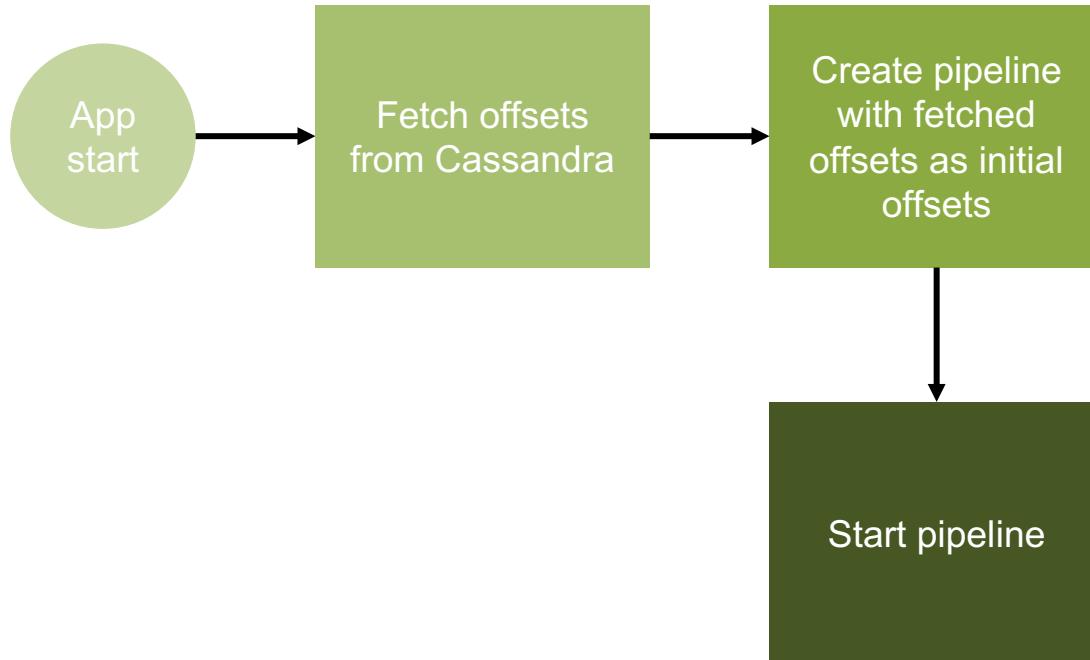
```
{  
  'key_1': [  
    {'event': '...', 'partition': 0, 'offset': 130},  
    {'event': '...', 'partition': 1, 'offset': 230},  
  ],  
  'key_2': [  
    {'event': '...', 'partition': 0, 'offset': 119},  
    {'event': '...', 'partition': 1, 'offset': 251},  
  ],  
}
```



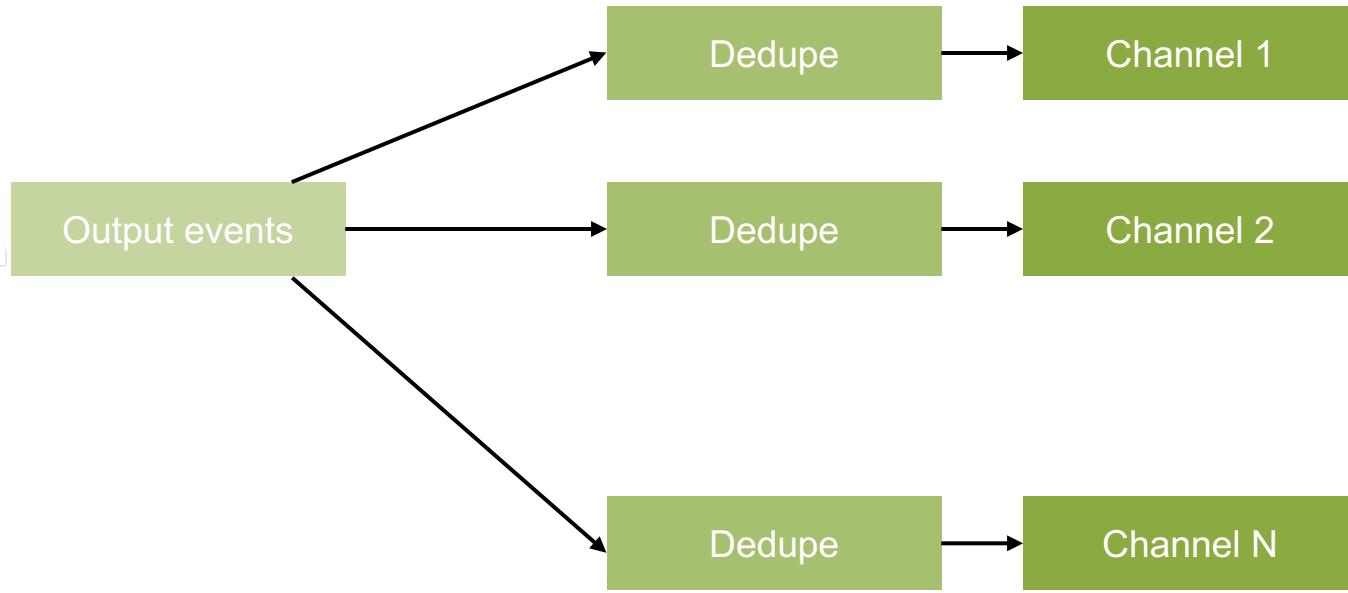
Find min offset per
cluster / topic / partition



Recovery via replay



Dealing with duplicates



Downsides



- Need to handle large data volume after deployments
- Size of state is limited by ability to playback



Looking forward...

- Possibly externalize state entirely
- Structured streaming may fix this... eventually



Conclusion



- Use `mapWithState` / `updateStateByKey` to deal with out-of-order events or complex logic needs
- State is not carried over after code changes
- Structured streaming may provide a better solution in the future





We're Hiring!

www.yelp.com/careers/



[fb.com/YelpEngineers](https://www.facebook.com/YelpEngineers)



[@YelpEngineering](https://twitter.com/YelpEngineering)



engineeringblog.yelp.com



github.com/yelp





Thank you!

- Yelp Ad Targeting at Scale with Apache Spark
 - Tomorrow (June 7th) @ 11.00 am
- Visit us in the expo hall at the **yelp** booth (T3)

