



Cassandra and SparkSQL

You don't need Functional Programming for Fun!

Russell (left) and Cara (right)

- Software Engineer
-  DATASTAX
- Spark-Cassandra Integration since Spark 0.9
 - Cassandra since Cassandra 1.2
 - 2 Year Scala Convert
 - Still not comfortable talking about Monads in public



A Story in 3 Parts

- Why SparkSQL?
- The Spark SQL Thrift Server
- Writing SQL for Spark

You have lots of options why Spark SQL?

- Scala?
- Java?
- Python?
- R?
- Notebooks?
- Erlang?



Spark is A Powerful Analytics Tool Built on Scala

Distributed Analytics Platform
with In Memory Capabilities

Lots of new concepts:

- RDDs
- DataSets
- Streaming
- Serialization
- Functional Programming



Functional Programming Is Awesome

Type Matching

Scala

Easy Parallelization

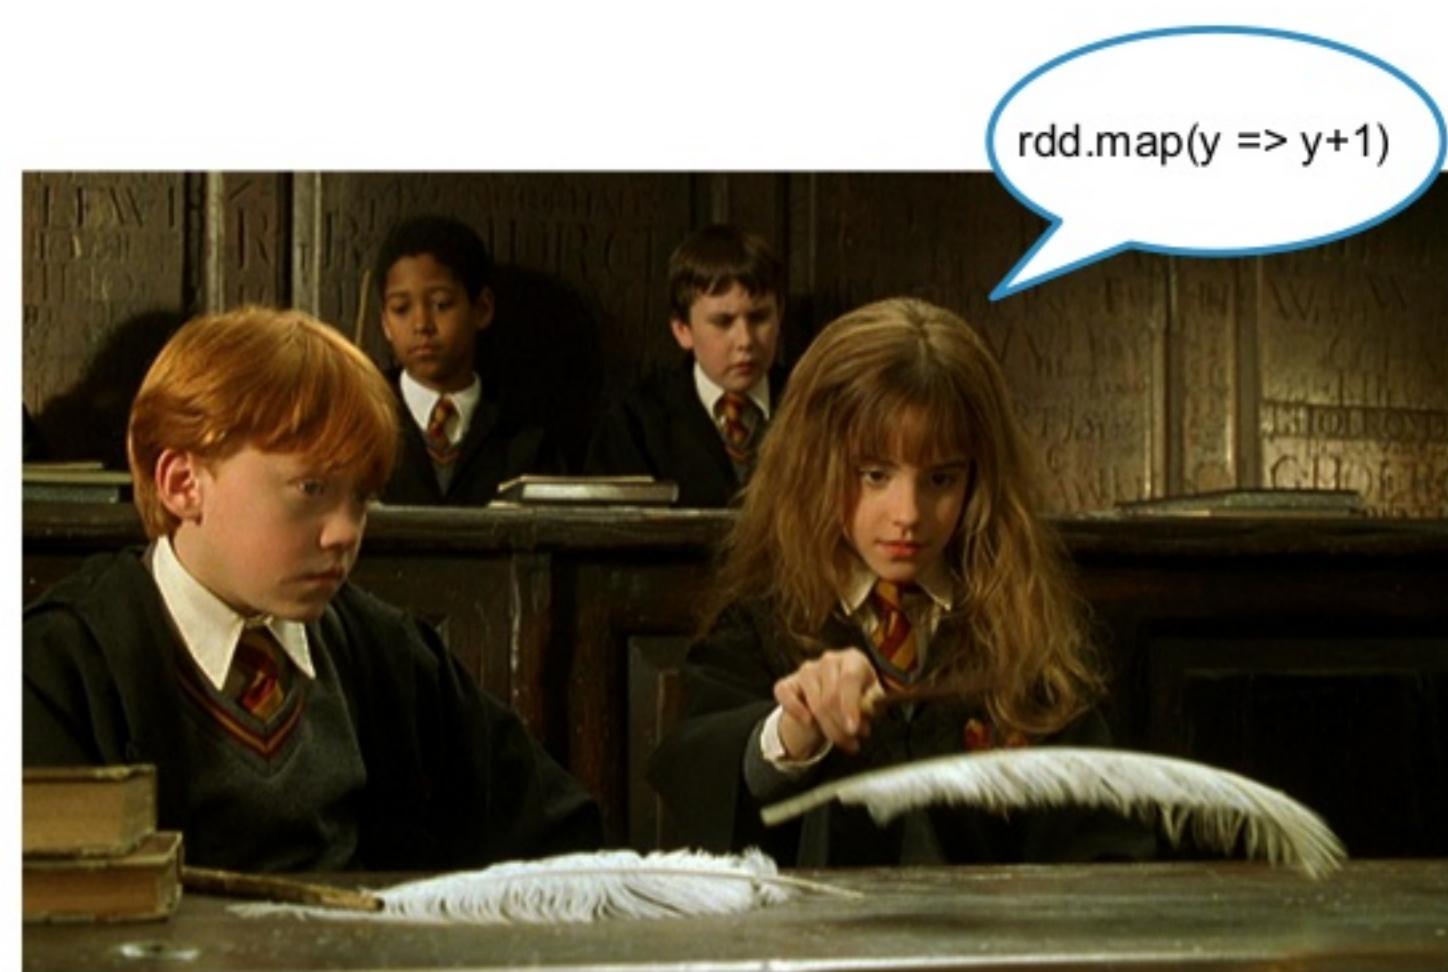
Anonymous Functions

Side-effect Free Functions

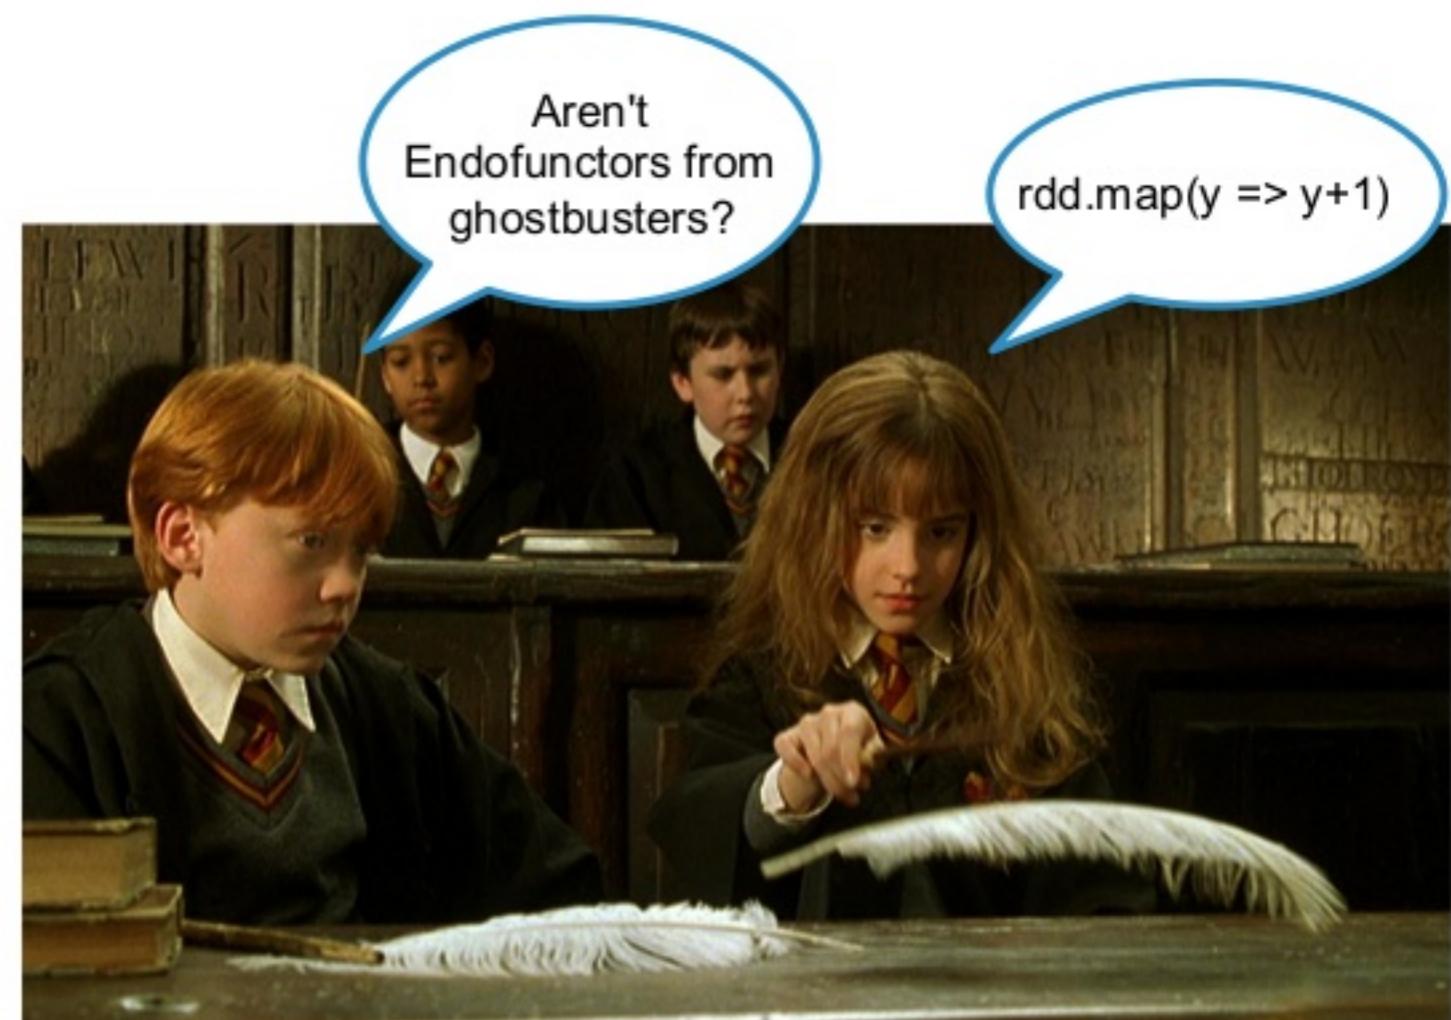
Monads

Endofunctors

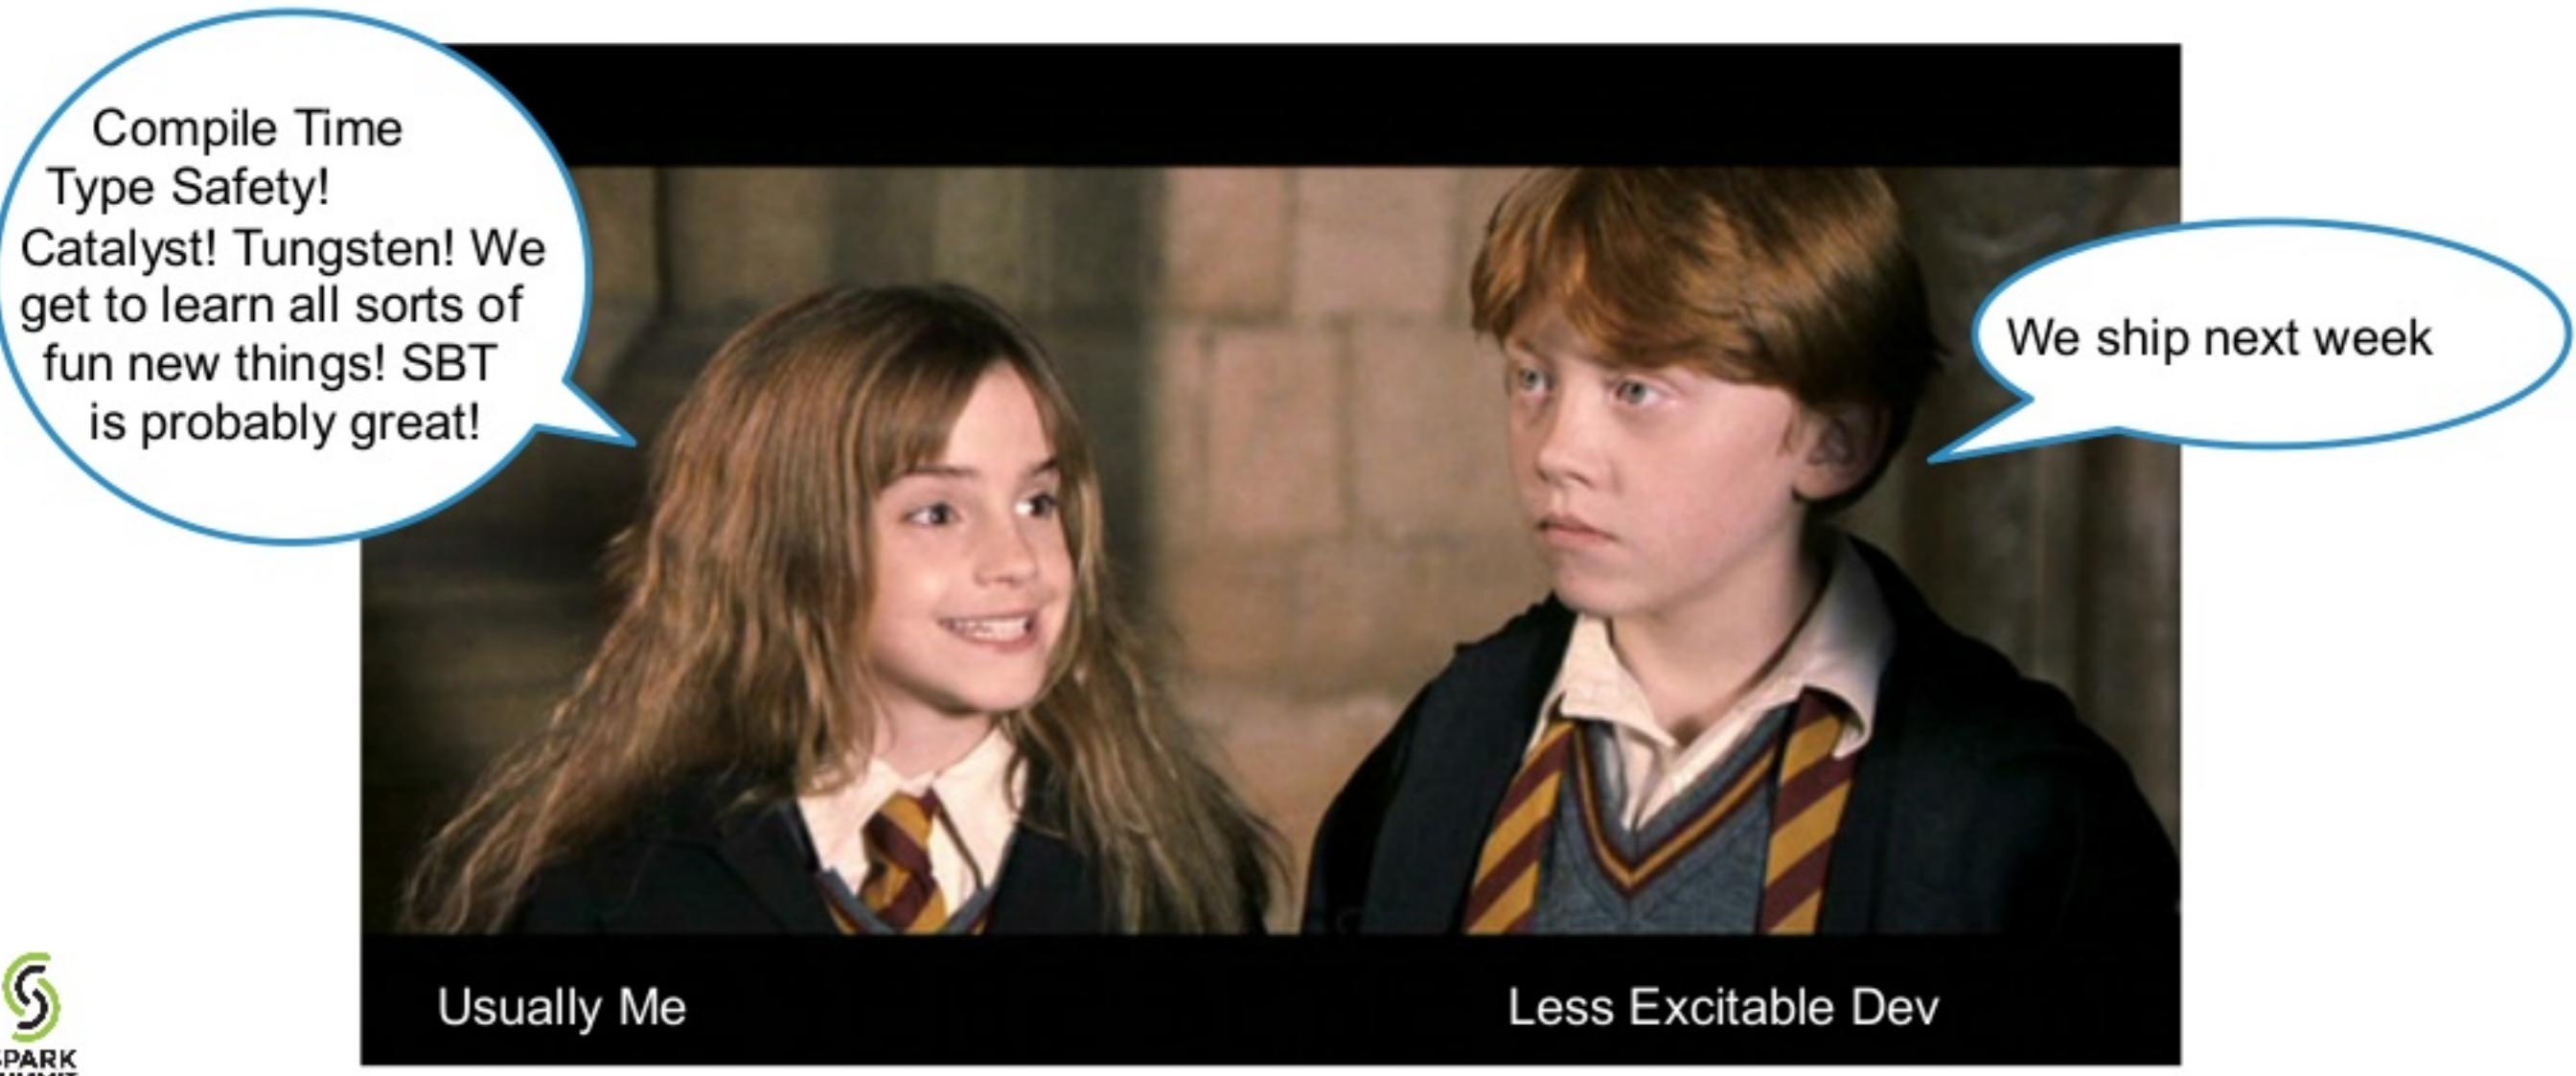
Async Models



Functional Programming can be Hard



Learning Takes Time



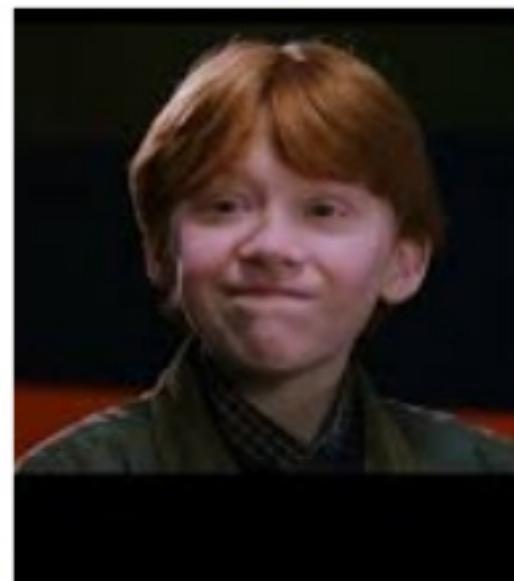
Spark SQL Provides A Familiar and Easy API

Use SQL to access the Power of Spark



Spark Sql Provides A Familiar and Easy API

SQL



Codegen!
Optimization!
Predicate Pushdowns



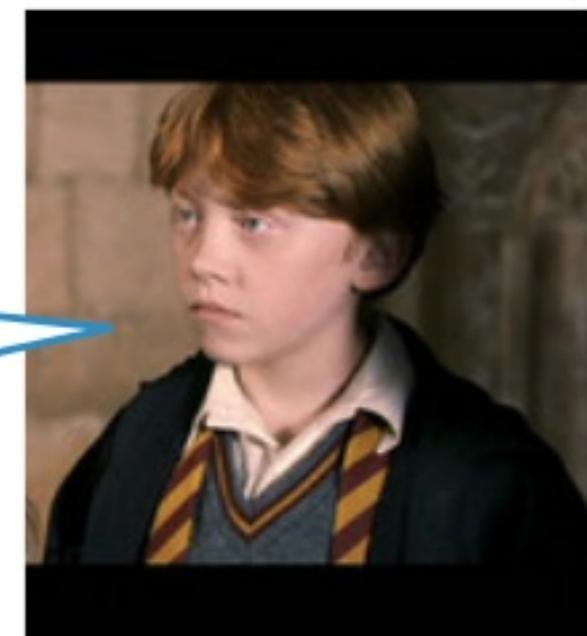
Catalyst

Distributed
Work

It still takes Scala/Java/Python/... Code.

```
import org.apache.spark.sql.cassandra._  
val df = spark  
  .read  
  .cassandraFormat("tab", "ks")  
  .load  
  df.createTempView("tab")  
spark.sql("SELECT * FROM tab").show  
+---+---+---+  
| k| c| v|  
+---+---+---+  
| 1| 1| 1|  
| 1| 2| 2|
```

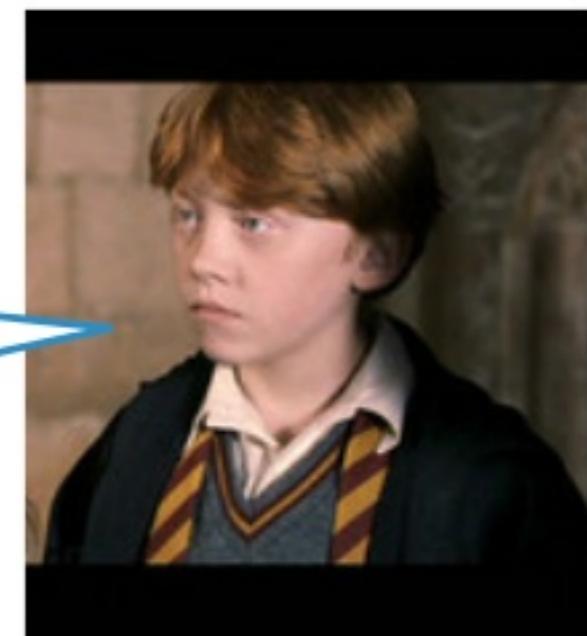
Let me color
code that by parts I
like vs parts I don't
like.



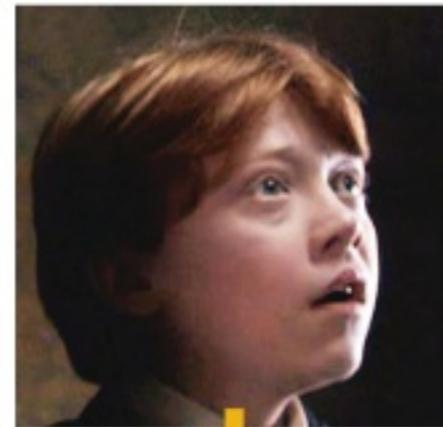
It still takes Scala/Java/Python/... Code.

```
import org.apache.spark.sql.cassandra._  
val df = spark  
  .read  
  .cassandraFormat("tab", "ks")  
  .load  
  df.createTempView("tab")  
spark.sql("SELECT * FROM tab").show  
+---+---+---+  
| k| c| v|  
+---+---+---+  
| 1| 1| 1|  
| 1| 2| 2|
```

Also, your import
has an underscore
in it..

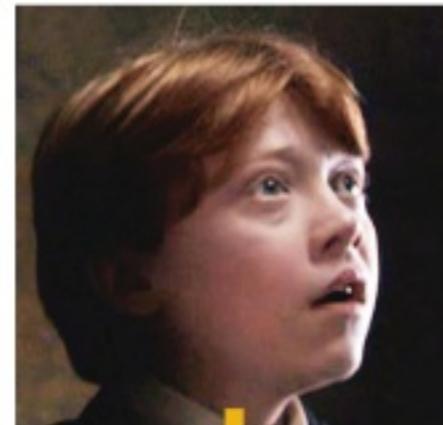


For exploration we have the Spark-SQL Shell



```
spark-sql> SELECT * FROM ks.tab;  
1    2    2  
1    3    3
```

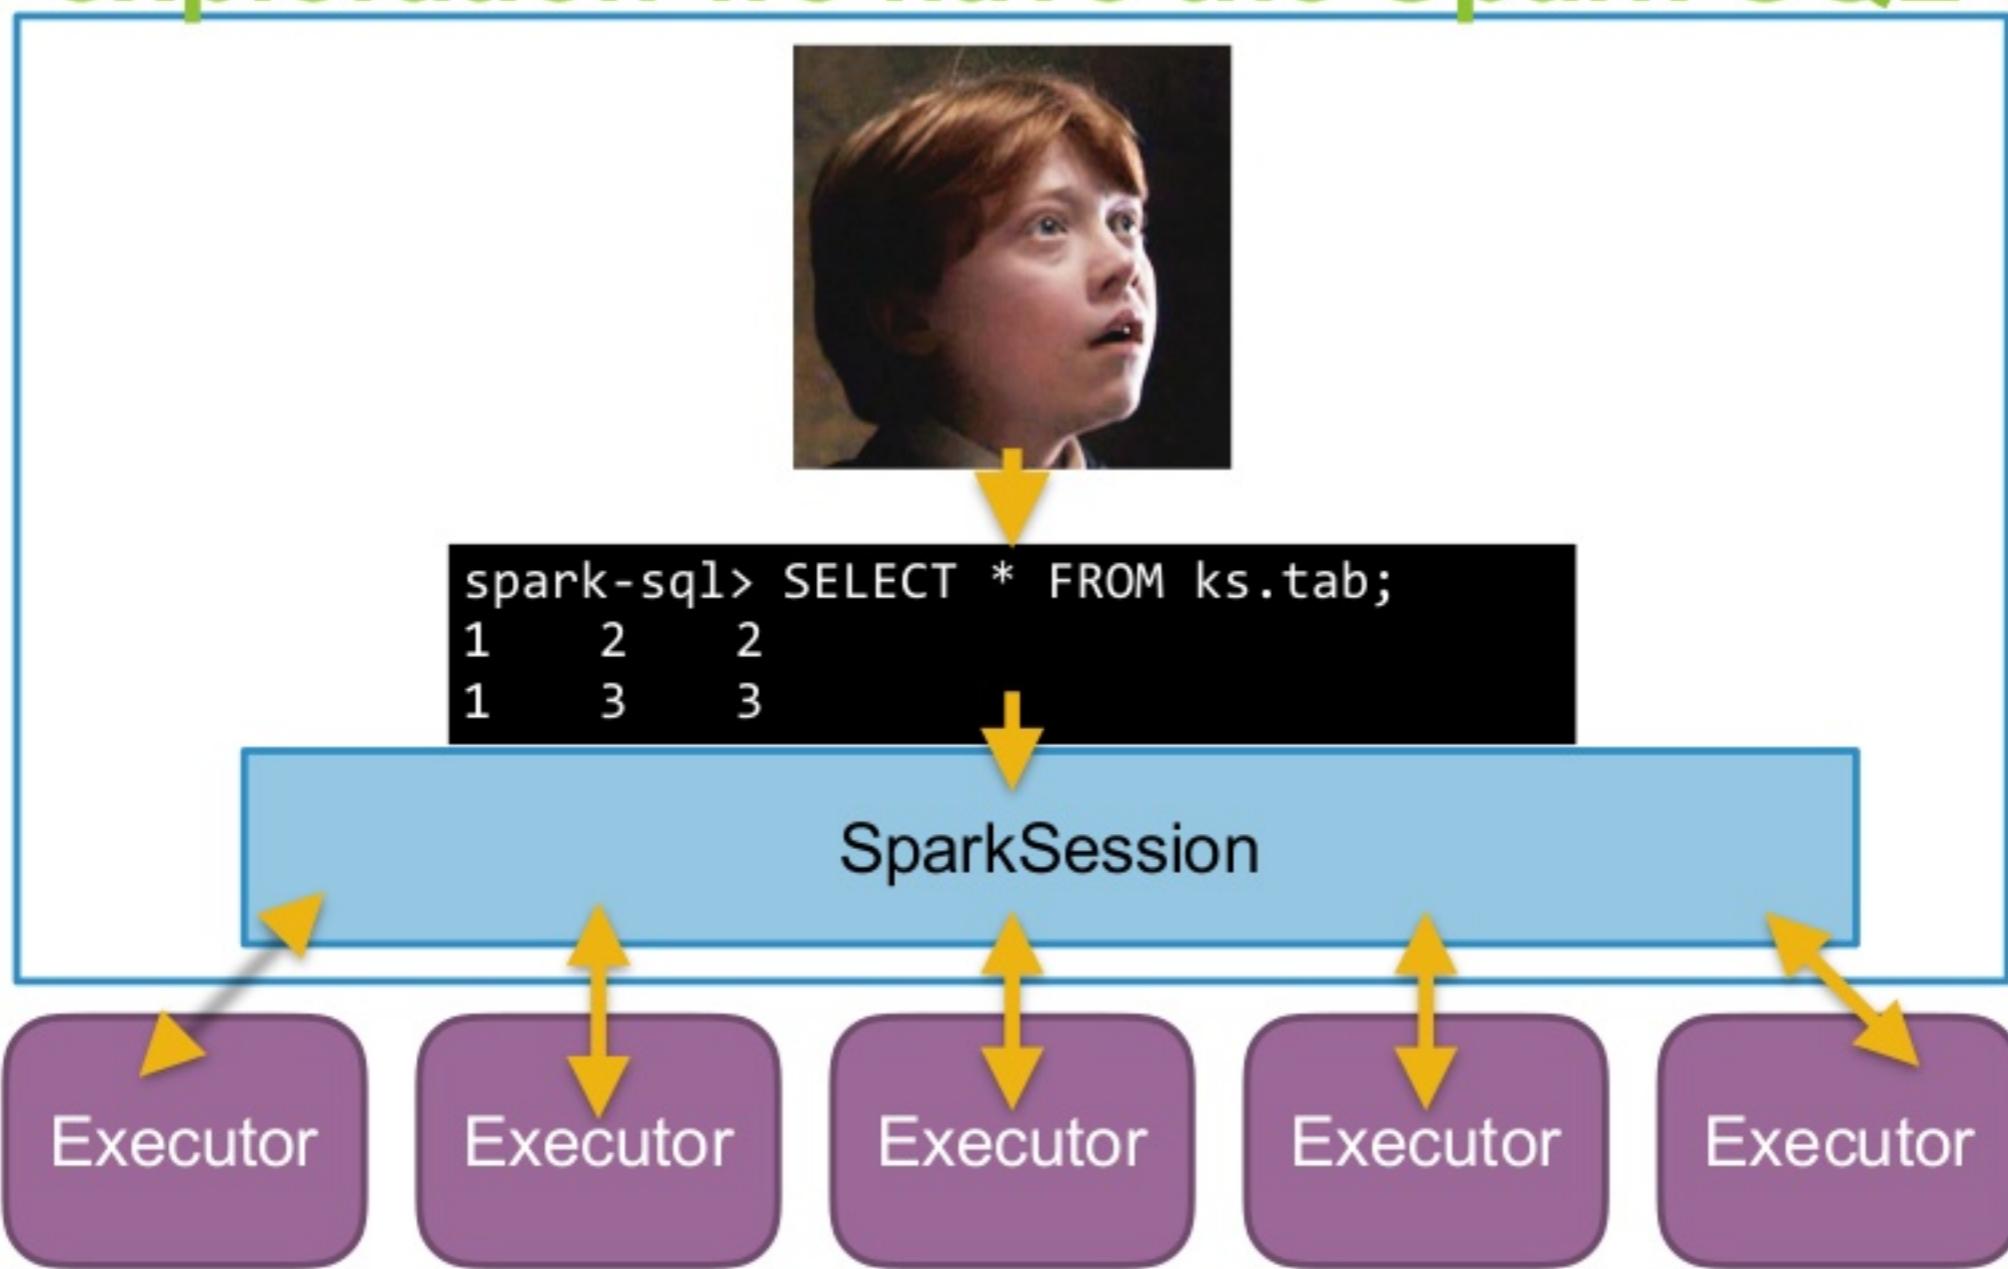
For exploration we have the Spark-SQL Shell



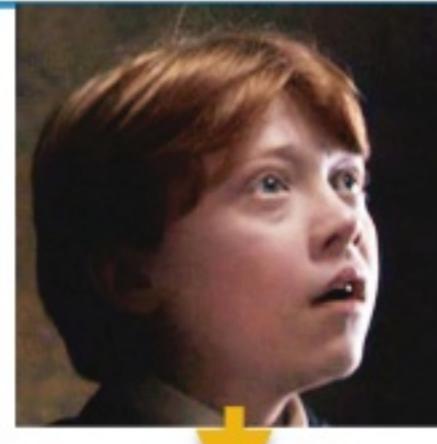
```
spark-sql> SELECT * FROM ks.tab;  
1    2    2  
1    3    3
```

SparkSession

For exploration we have the Spark-SQL Shell



Not really good for multiple-users



```
spark-sql> SELECT * FROM ks.tab;  
1    2    2  
1    3    3
```

SparkSession

Executor

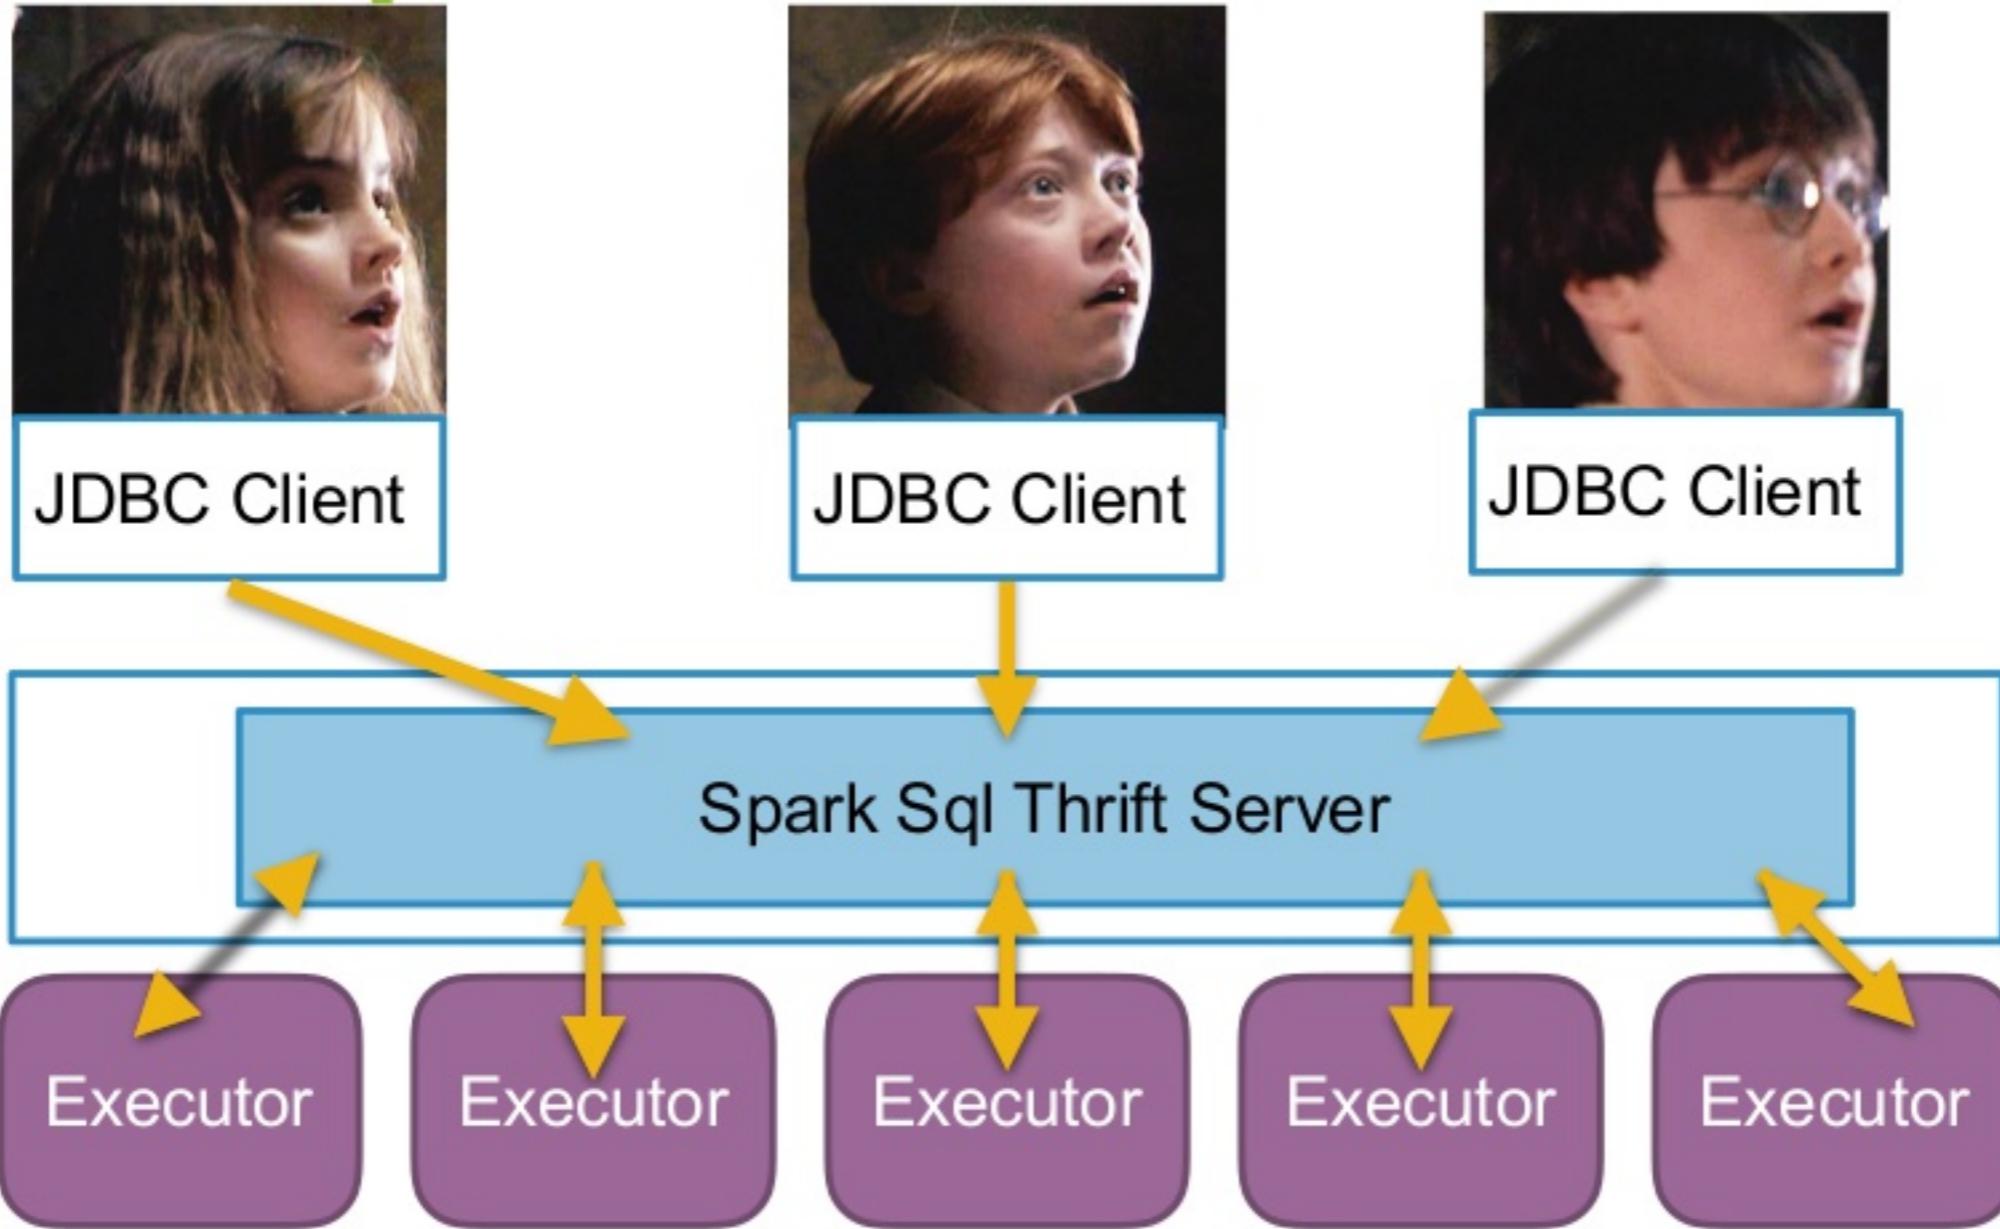
Executor

Executor

Executor

Executor

Enter Spark Thrift Server



The Spark Sql Thrift Server is a Spark Application

- Built on HiveServer2
- Single Spark Context
- Clients Communicate with it via JDBC
- Can use all SparkSQL
- Fair Scheduling
- Clients can share Cached Resources
- Security

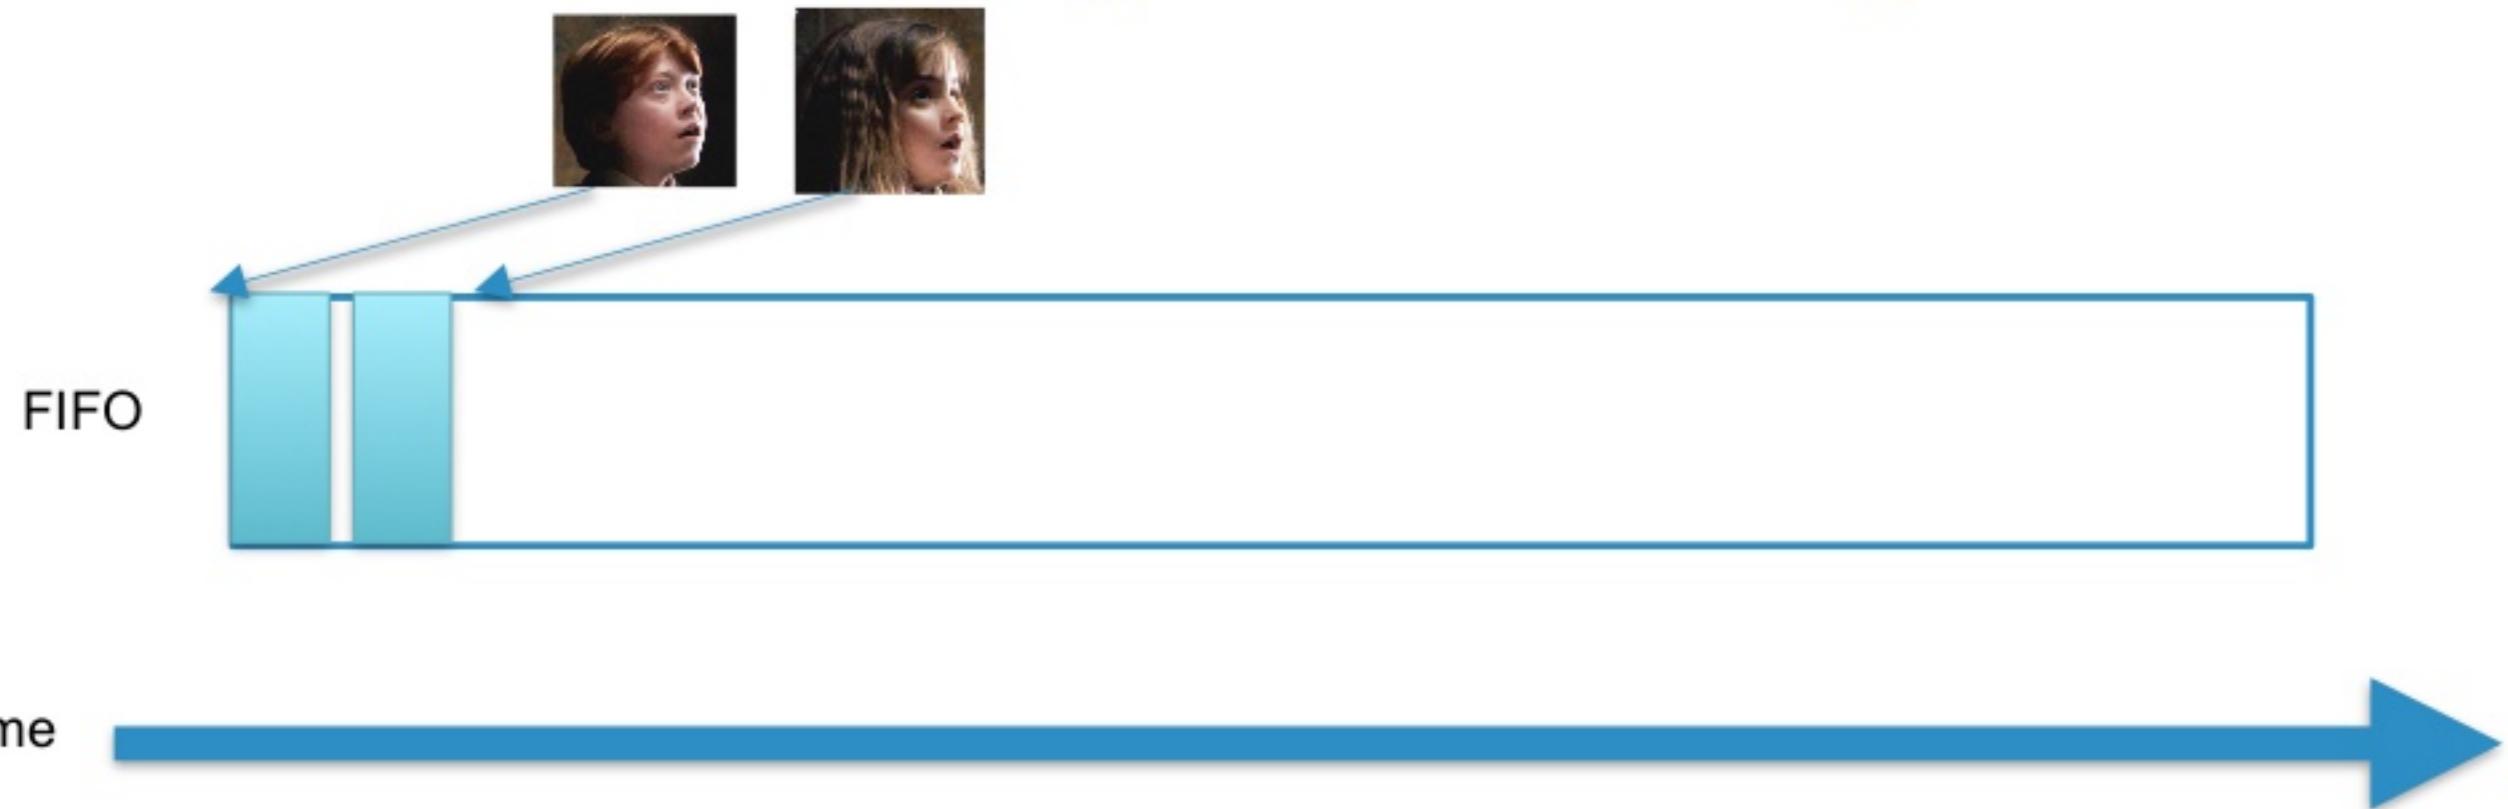
The Spark Sql ThriftServer is a Spark Application

- Built on HiveServer2
- Single Spark Context
- Clients Communicate with it via JDBC
- Can use all SparkSQL
- **Fair Scheduling**
- **Clients can share Cached Resources**
- Security

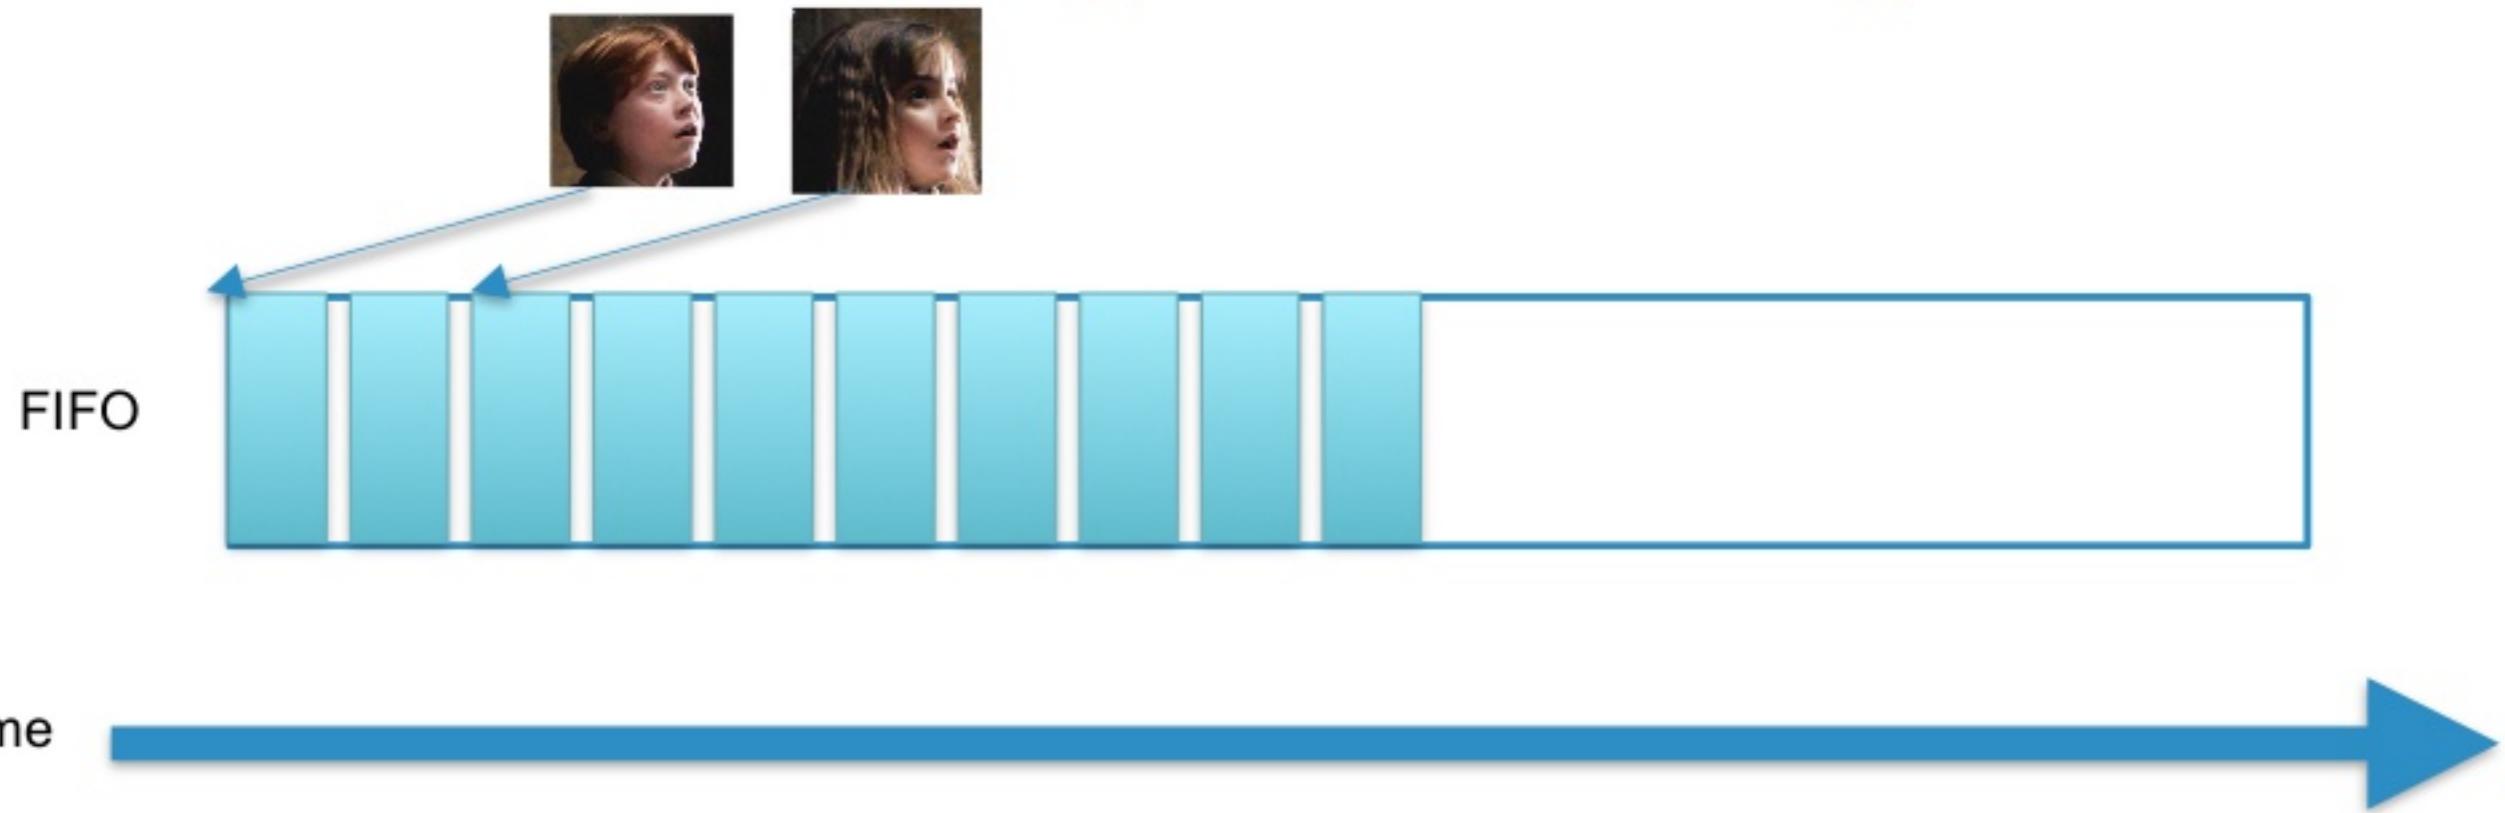
Fair Scheduling is Sharing



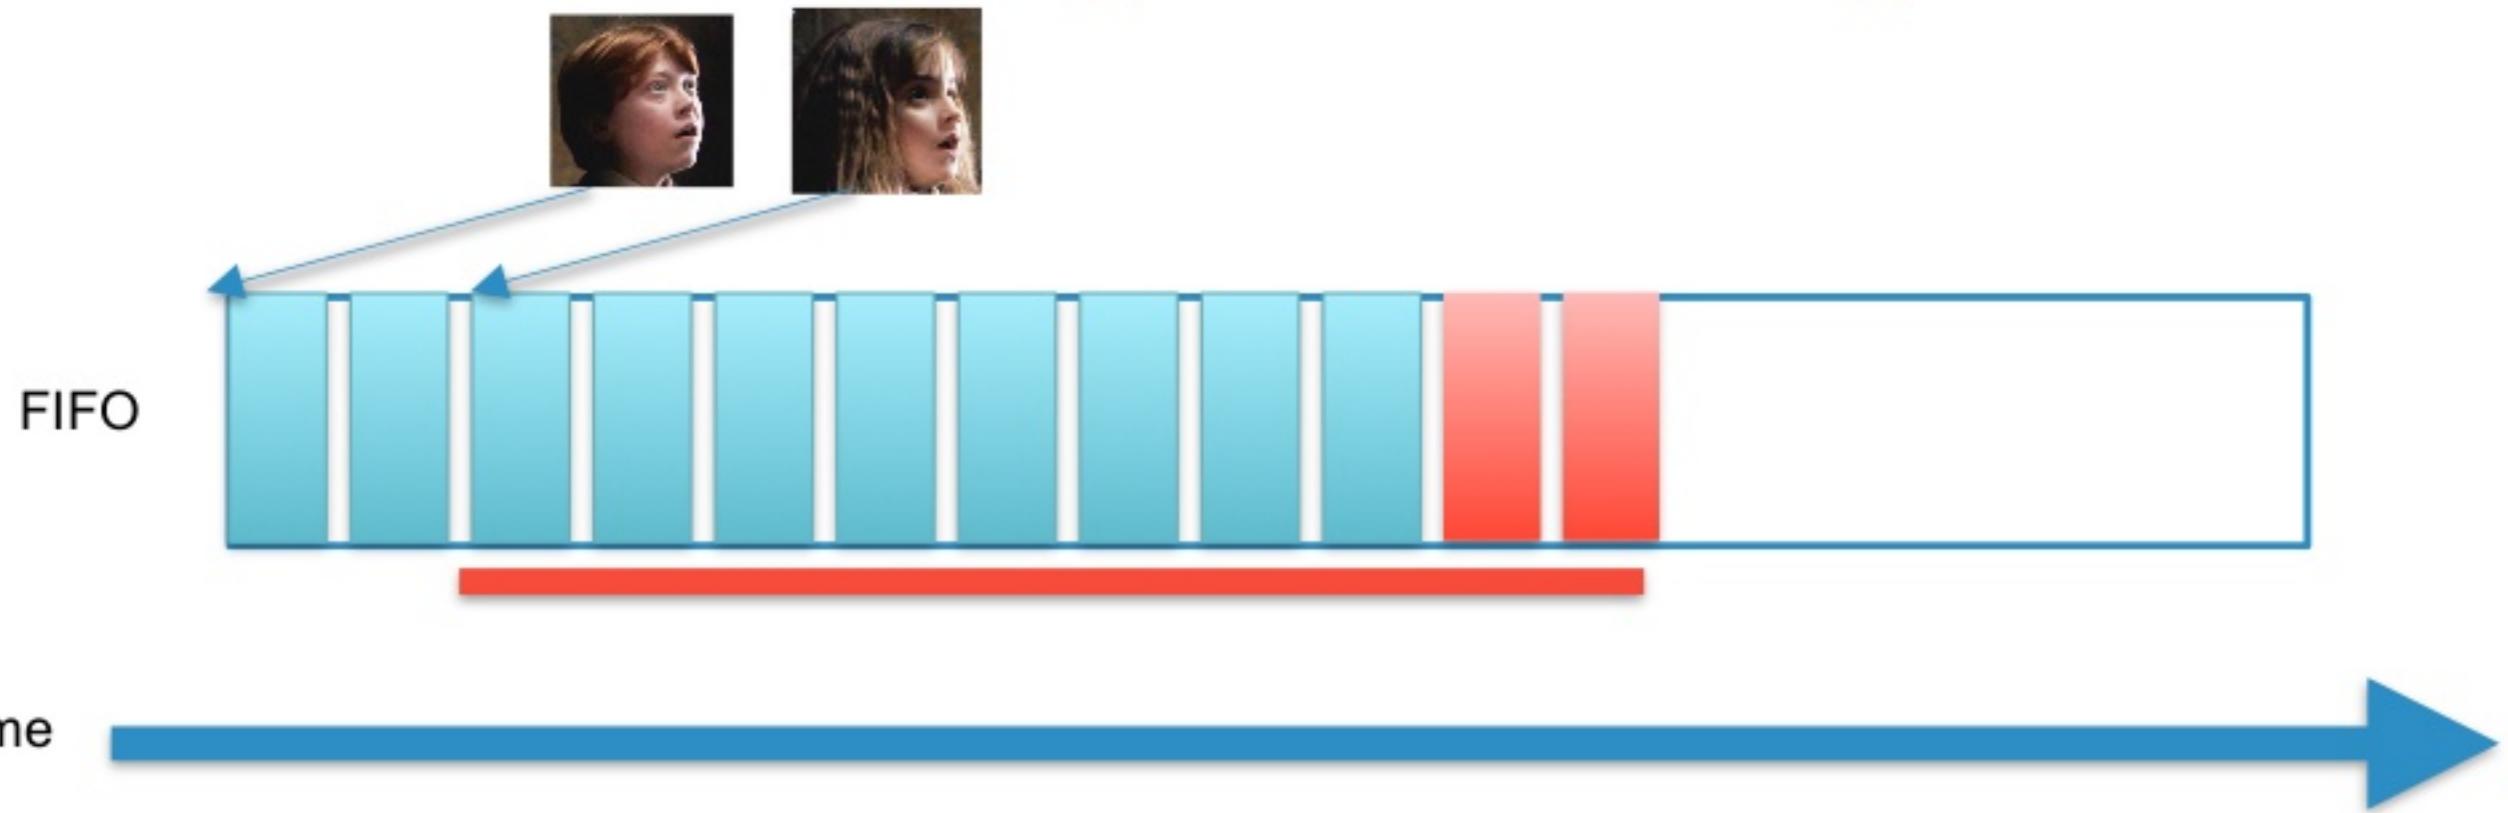
Fair Scheduling is Sharing



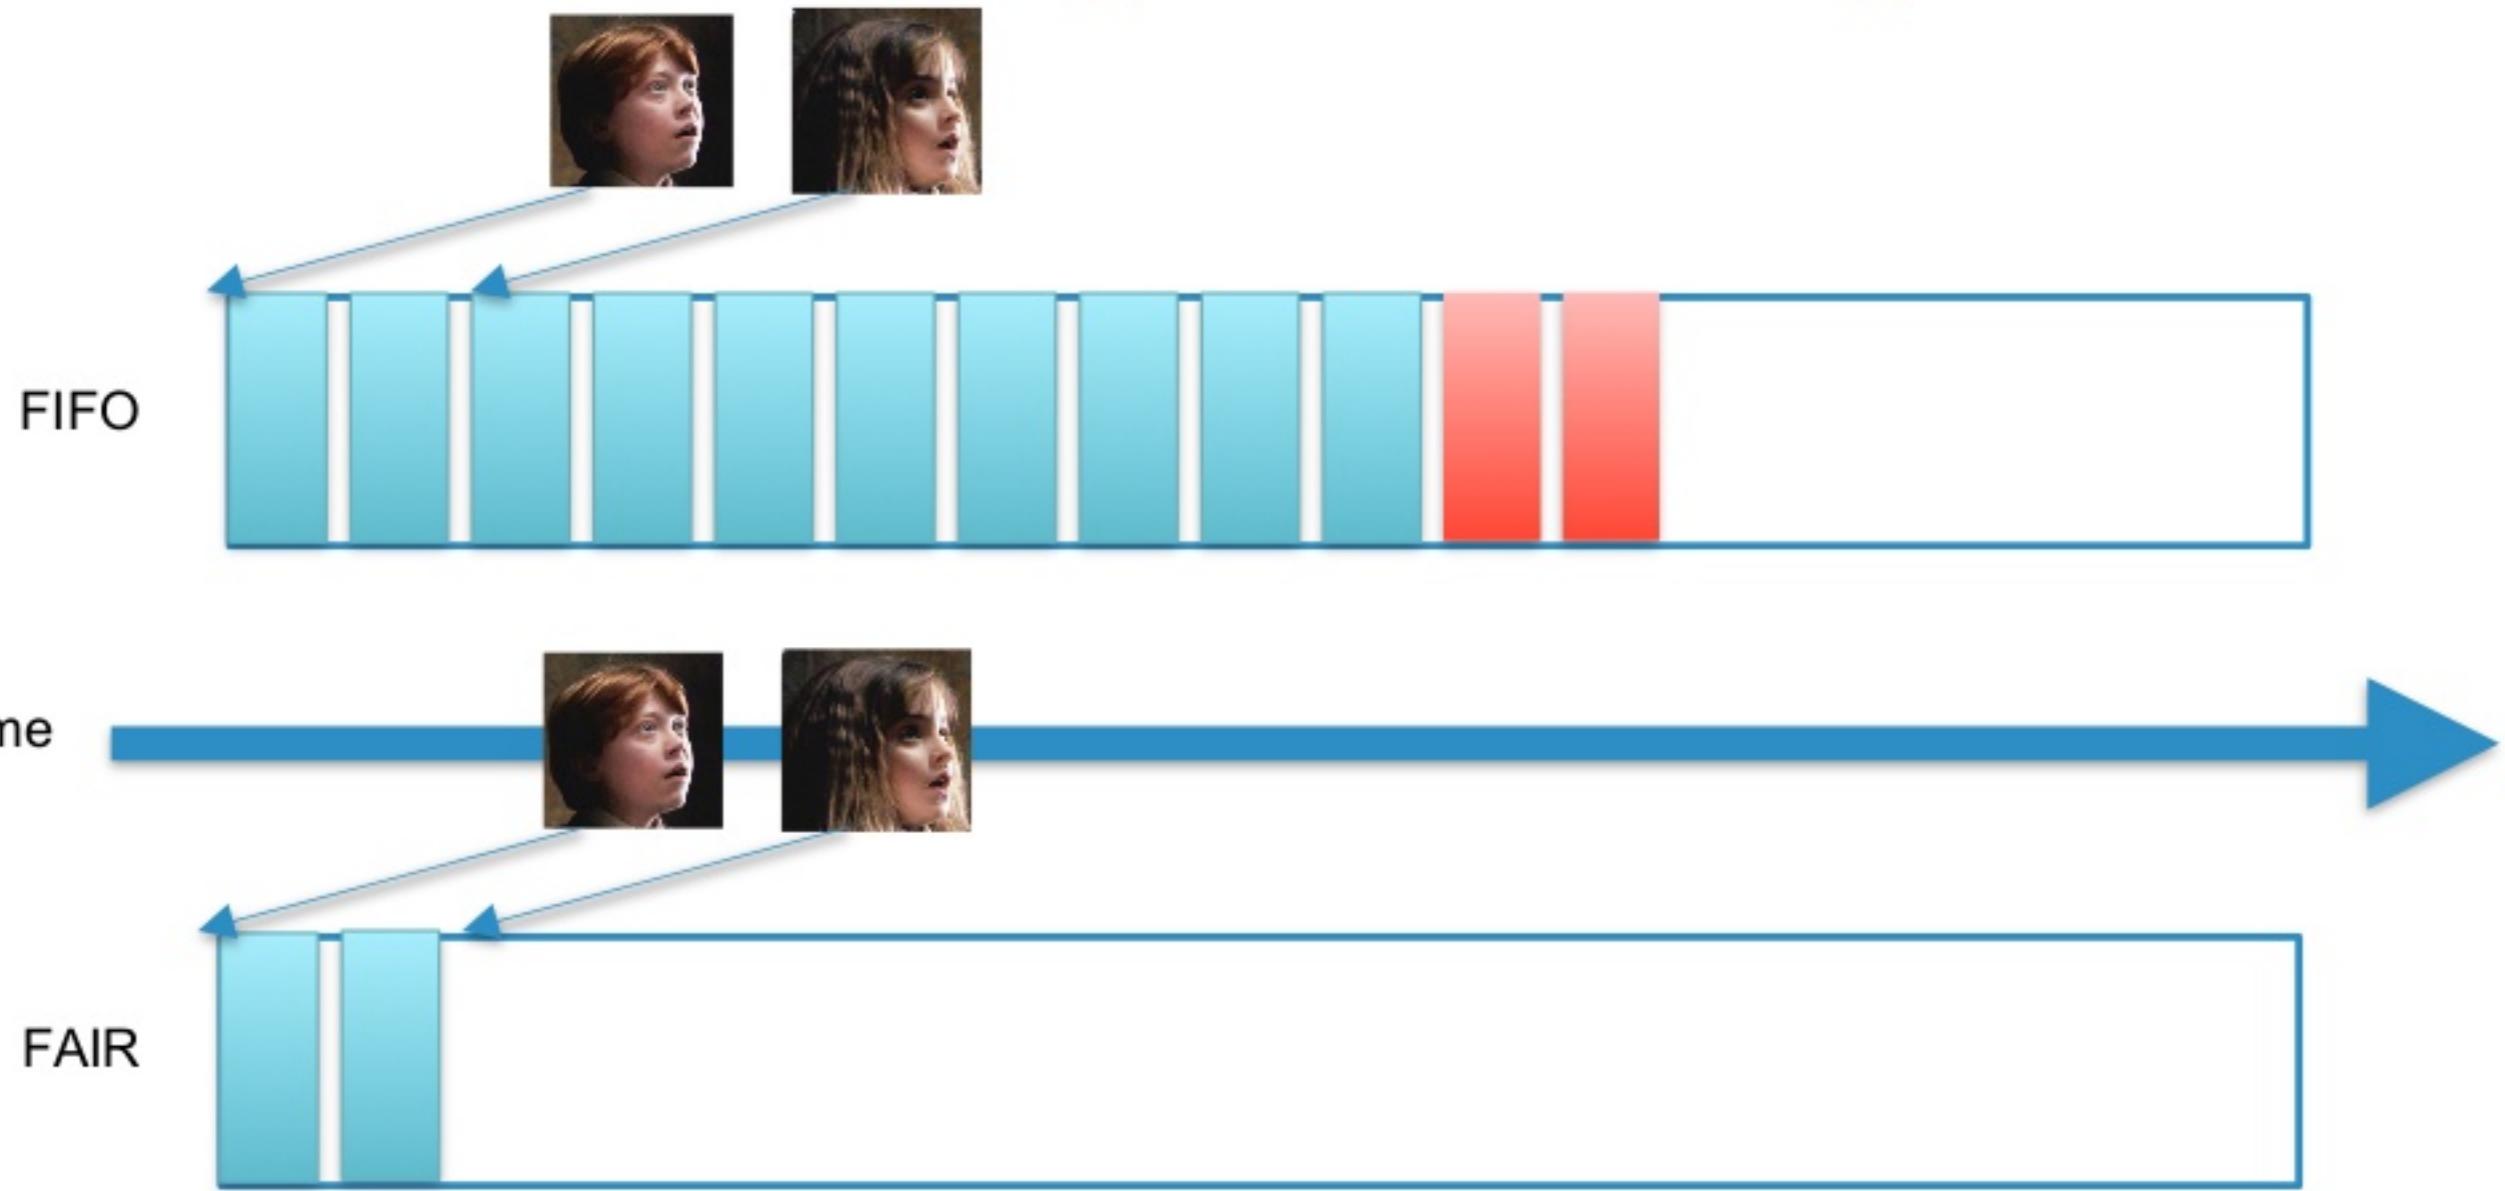
Fair Scheduling is Sharing



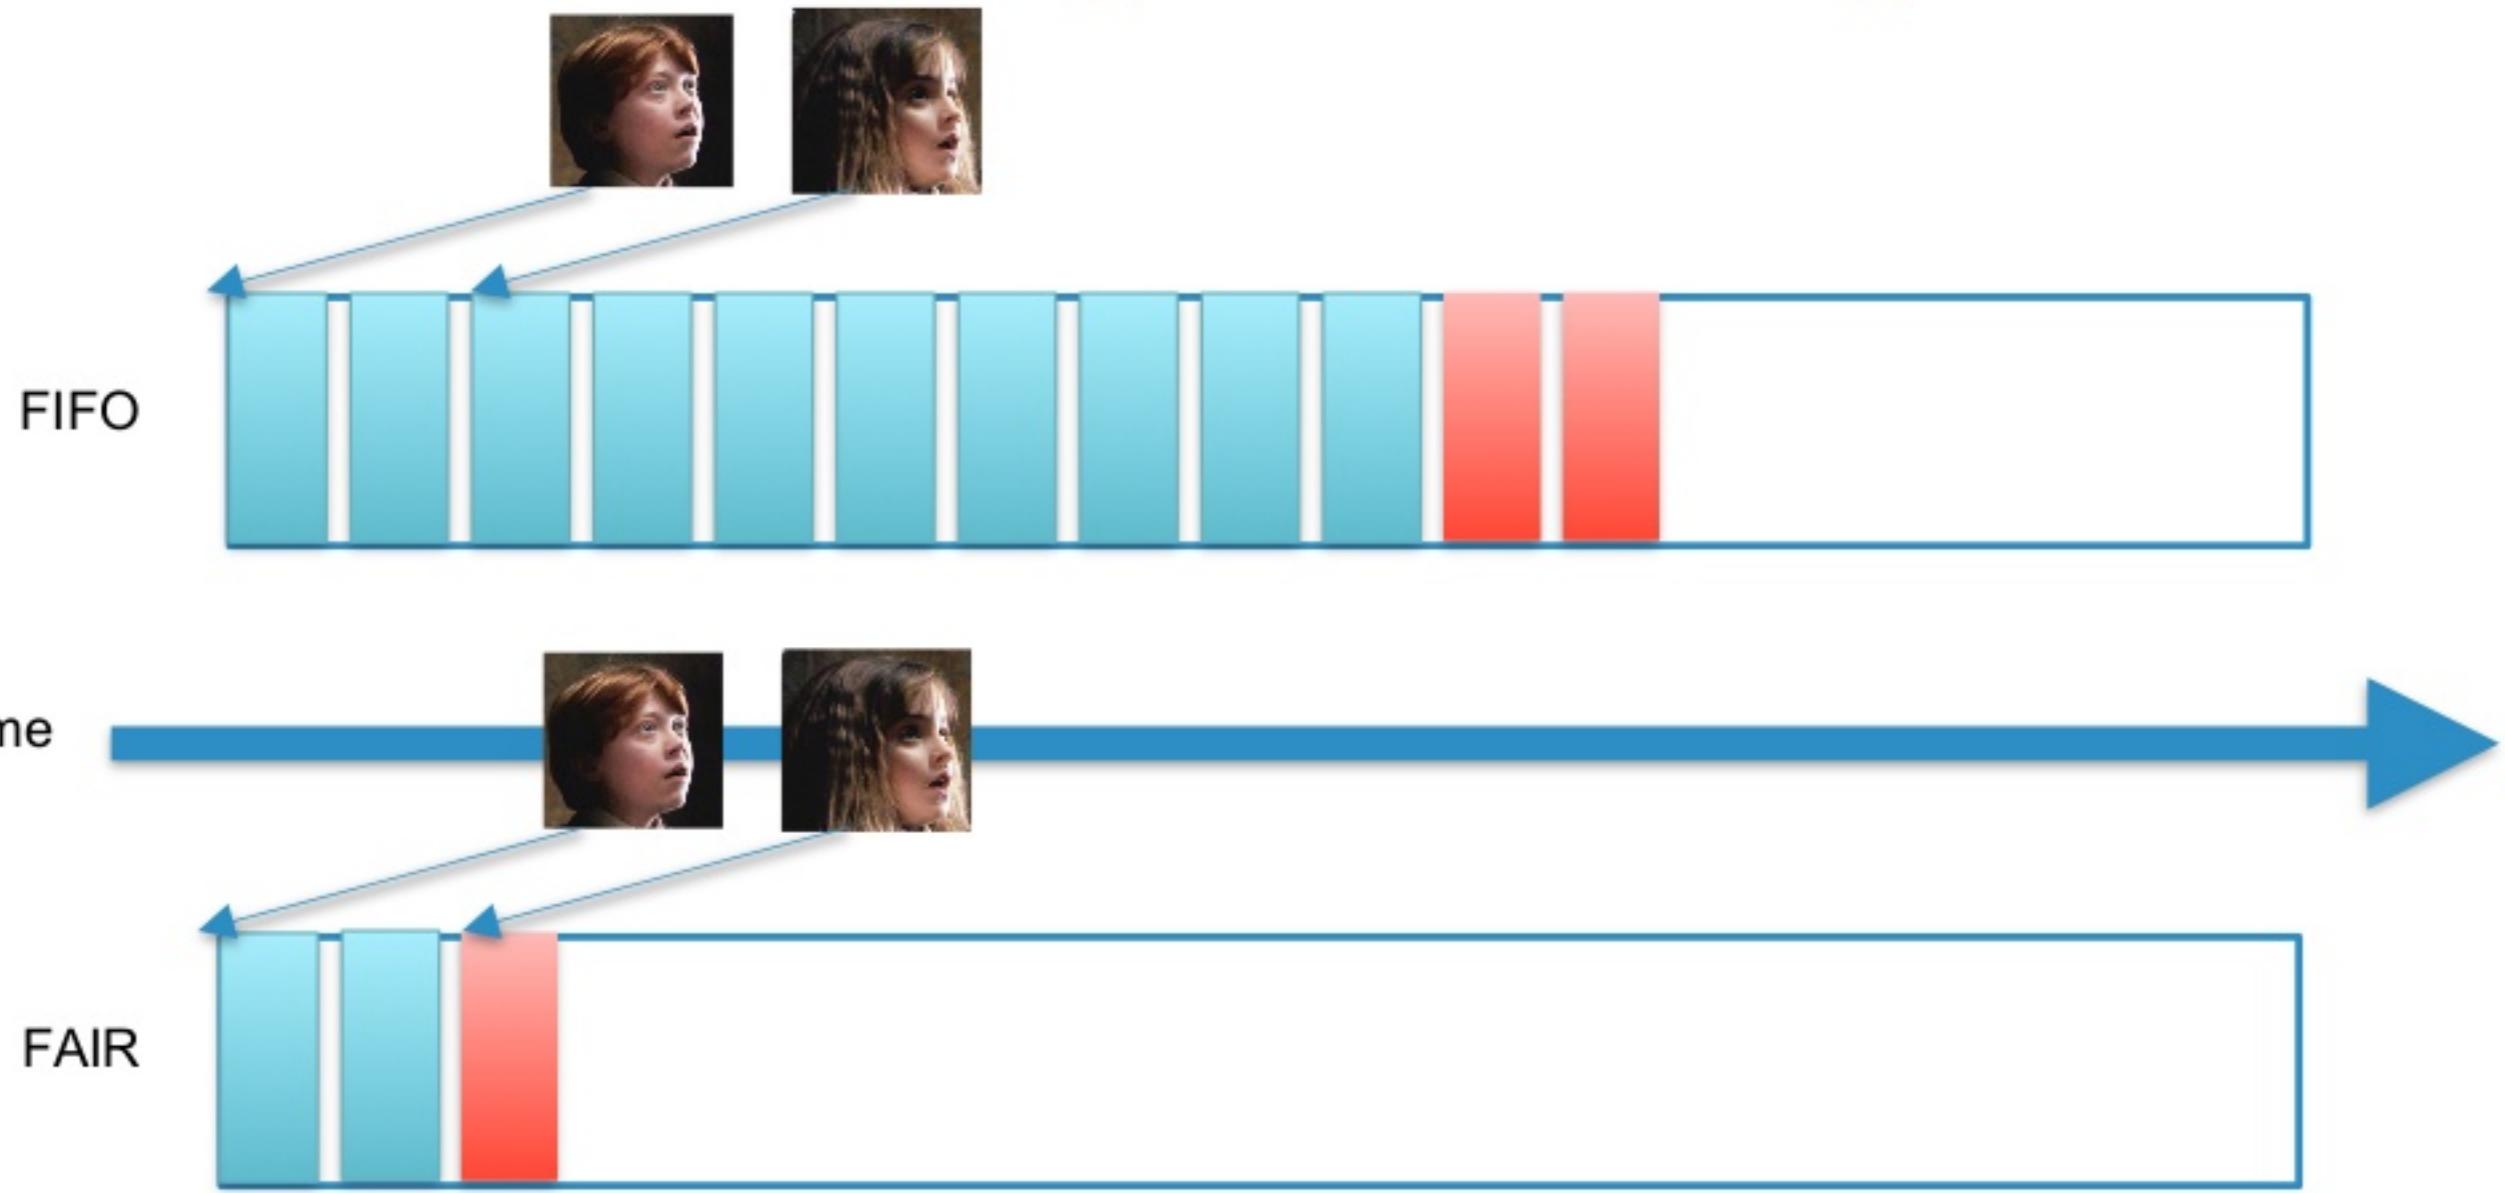
Fair Scheduling is Sharing



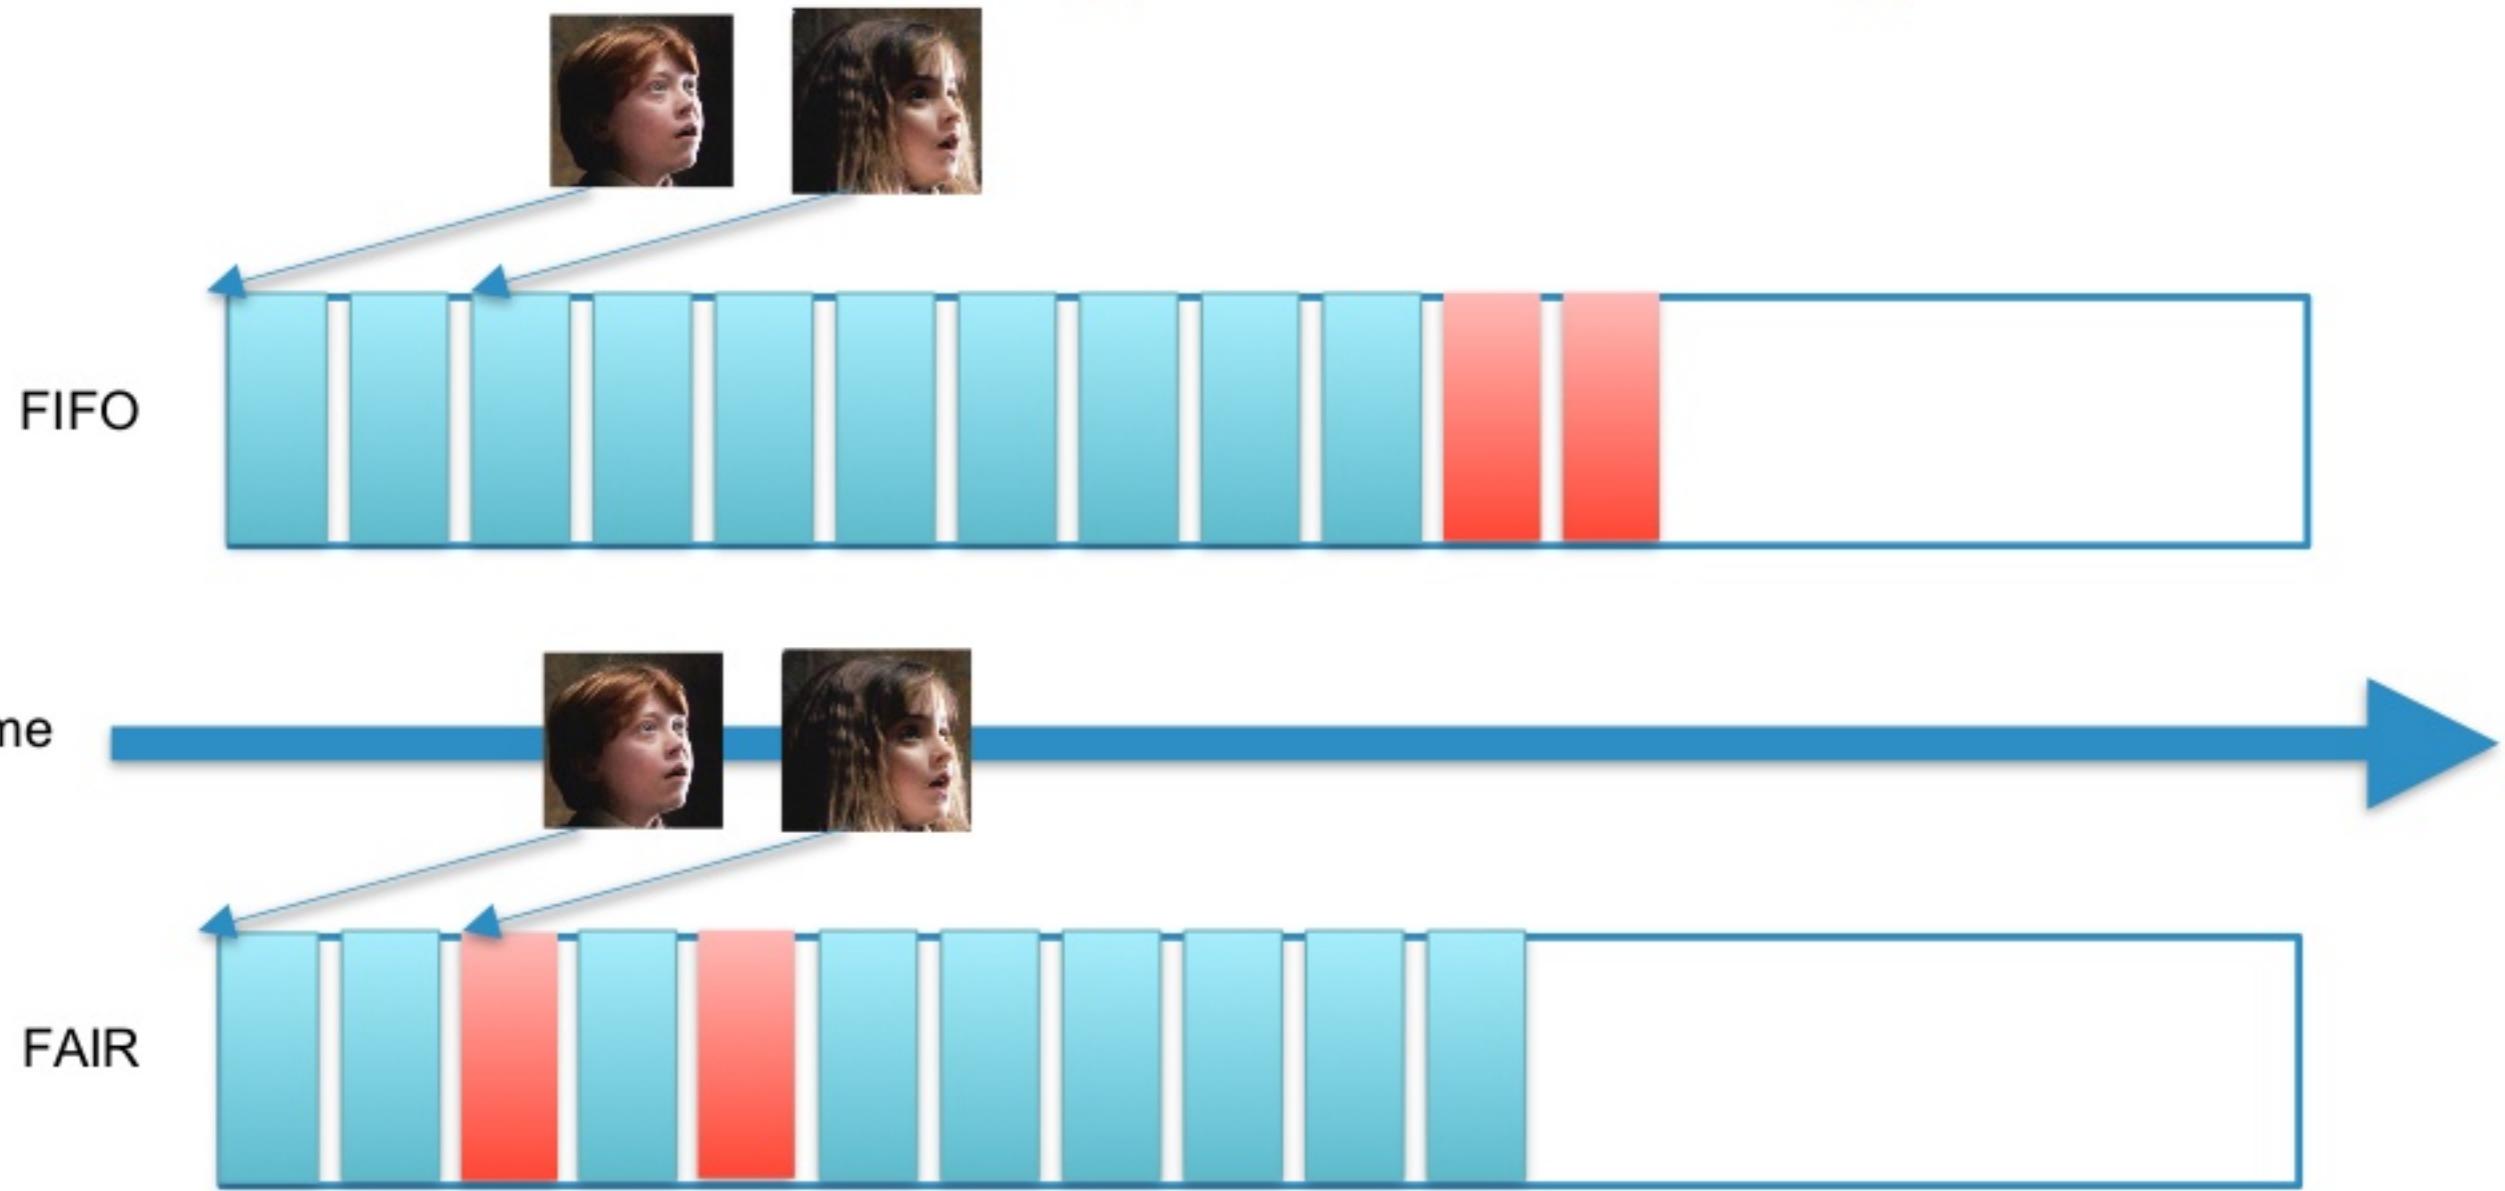
Fair Scheduling is Sharing



Fair Scheduling is Sharing

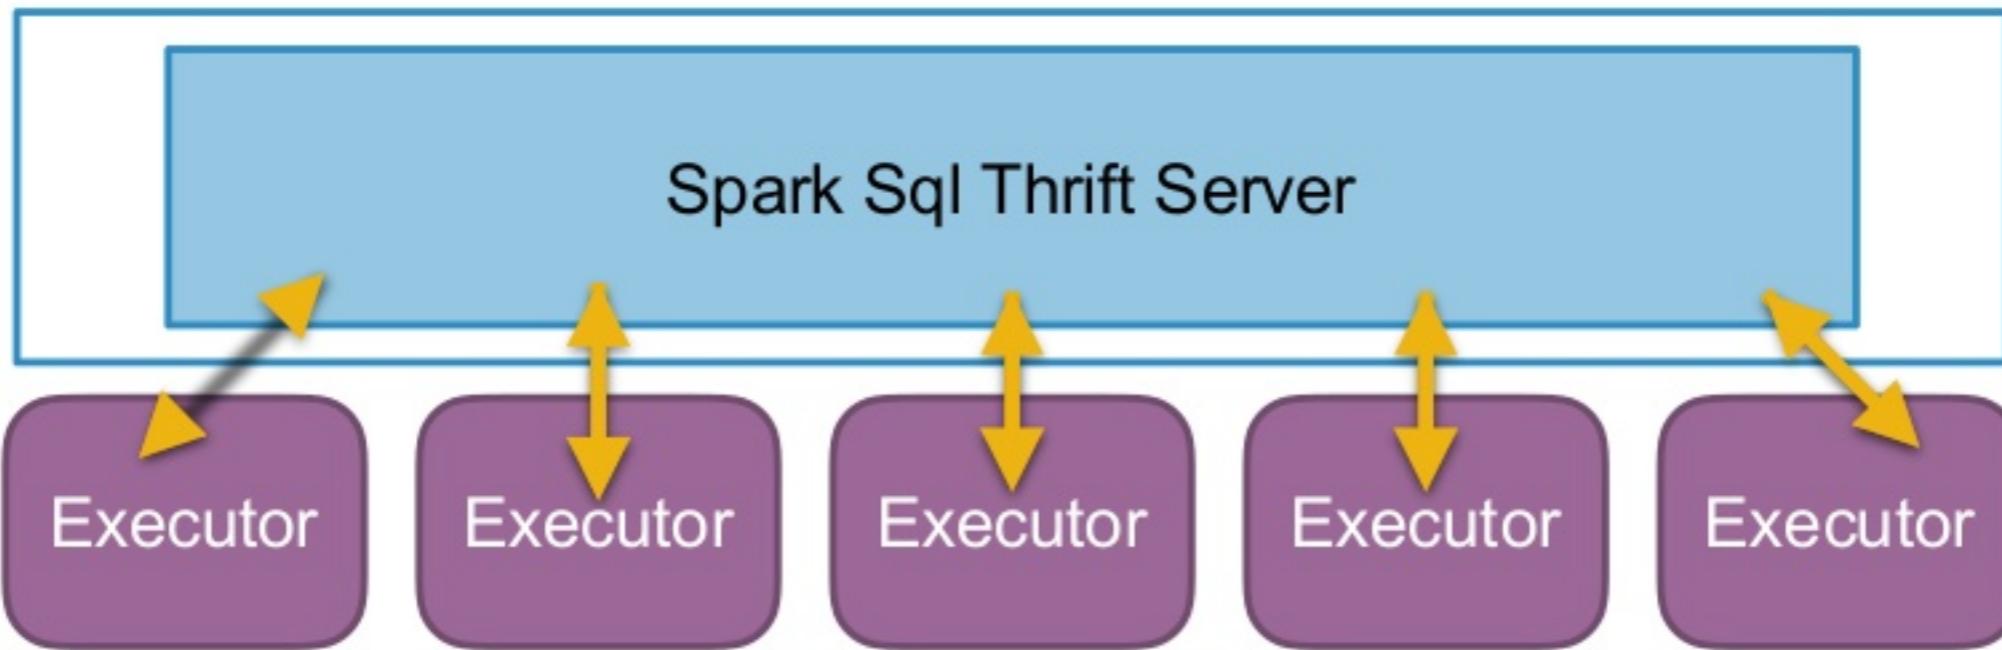


Fair Scheduling is Sharing



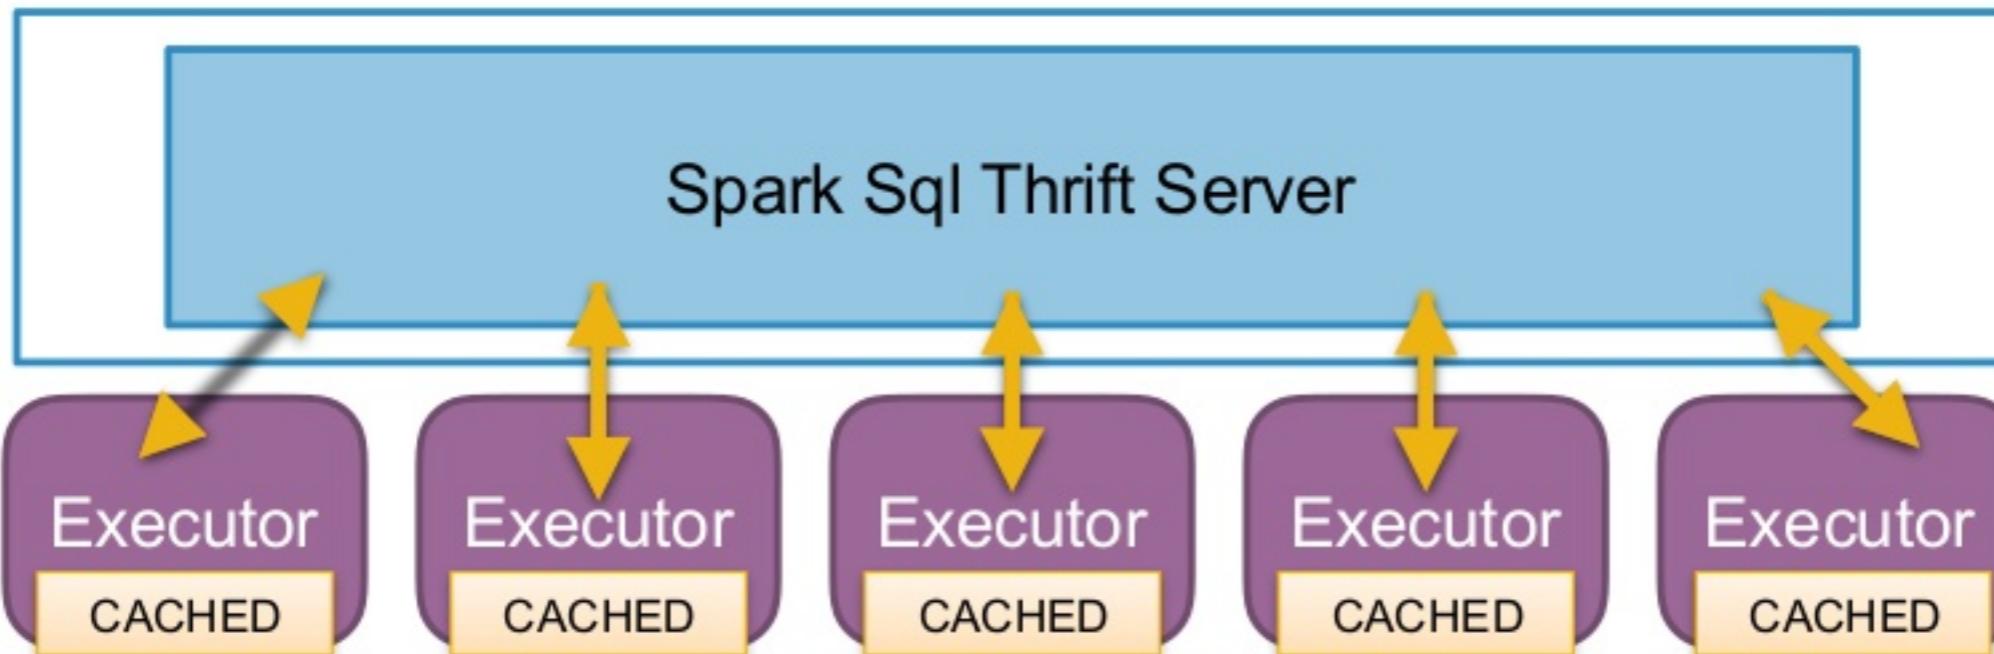
SingleContext can Share Cached Data

cache TABLE today select * from ks.tab where date = today;



SingleContext can Share Cached Data

cache TABLE today select * from ks.tab where date = today;

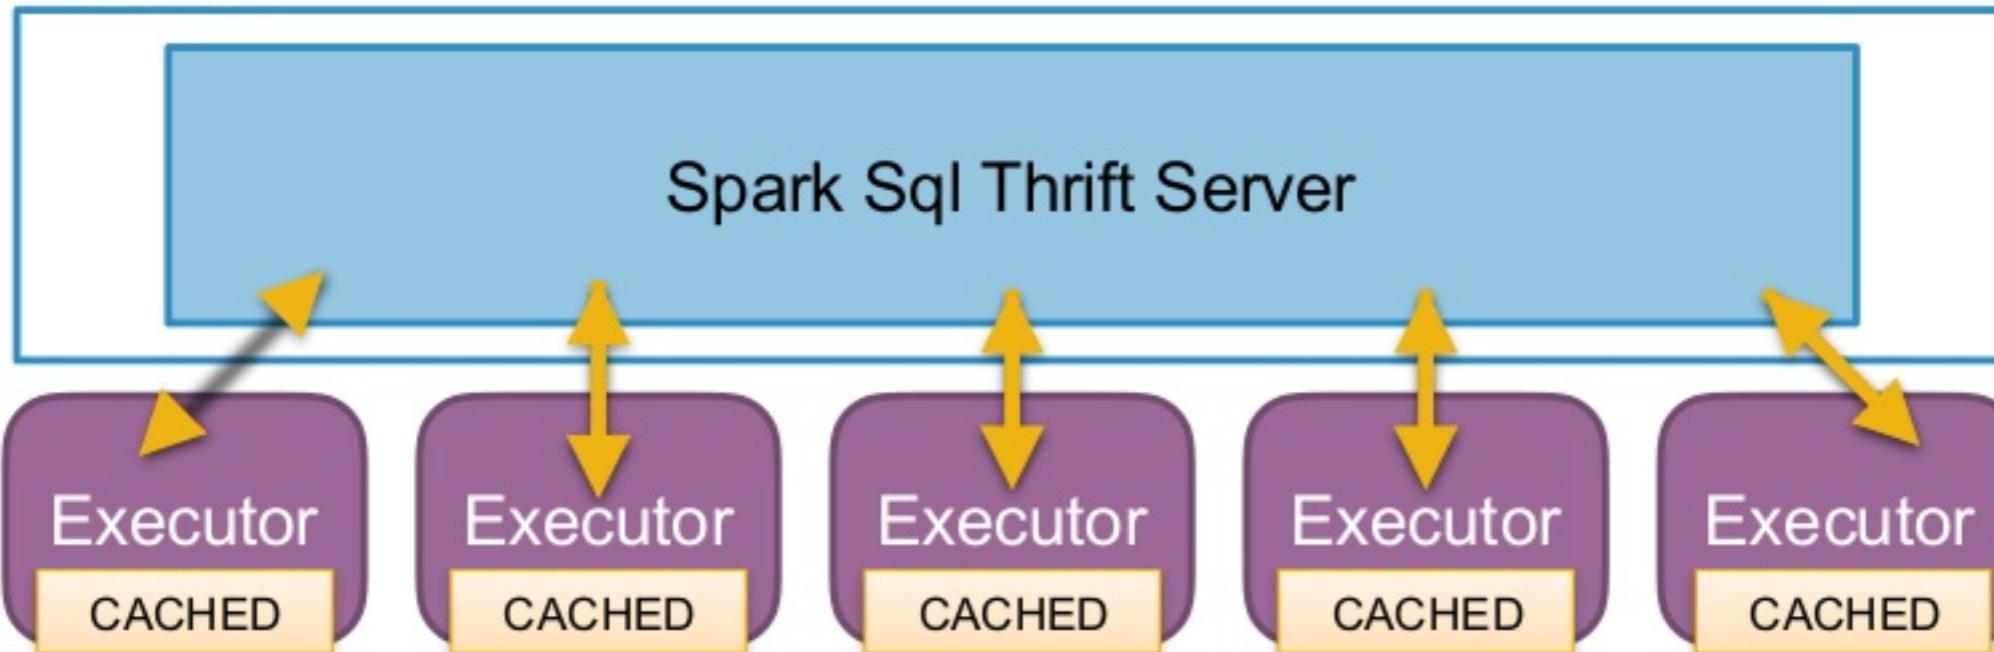


SingleContext can Share Cached Data

```
cache TABLE today select * from ks.tab where date = today;
```



```
SELECT * from TODAY where age > 5
```



How to use it

Starts from the command line and can use all Spark Submit Args

- ./sbin/start-thriftserver.sh
- dse spark-sql-thriftserver start

Use with all of your favorite Spark Packages like the Spark Cassandra Connector!

```
--packages com.datastax.spark:spark-cassandra-connector_2.11:2.0.2  
--conf spark.cassandra.connection.host=127.0.0.1
```

How to use it

Starts from the command line and can use all Spark Submit Args

- ./sbin/start-thriftserver.sh
- dse spark-sql-thriftserver start

```
starting org.apache.spark.sql.hive.thriftserver.HiveThriftServer2
```

Hive? Wait I thought we were Doing Spark

- Built on HiveServer2



Why does it say Hive everywhere?

```
starting org.apache.spark.sql.hive.thriftserver.HiveThriftServer2
```



A Brief History of the Spark Thrift Server

- Thrift?
- Hive?



They are not the Same



Time for a quick history

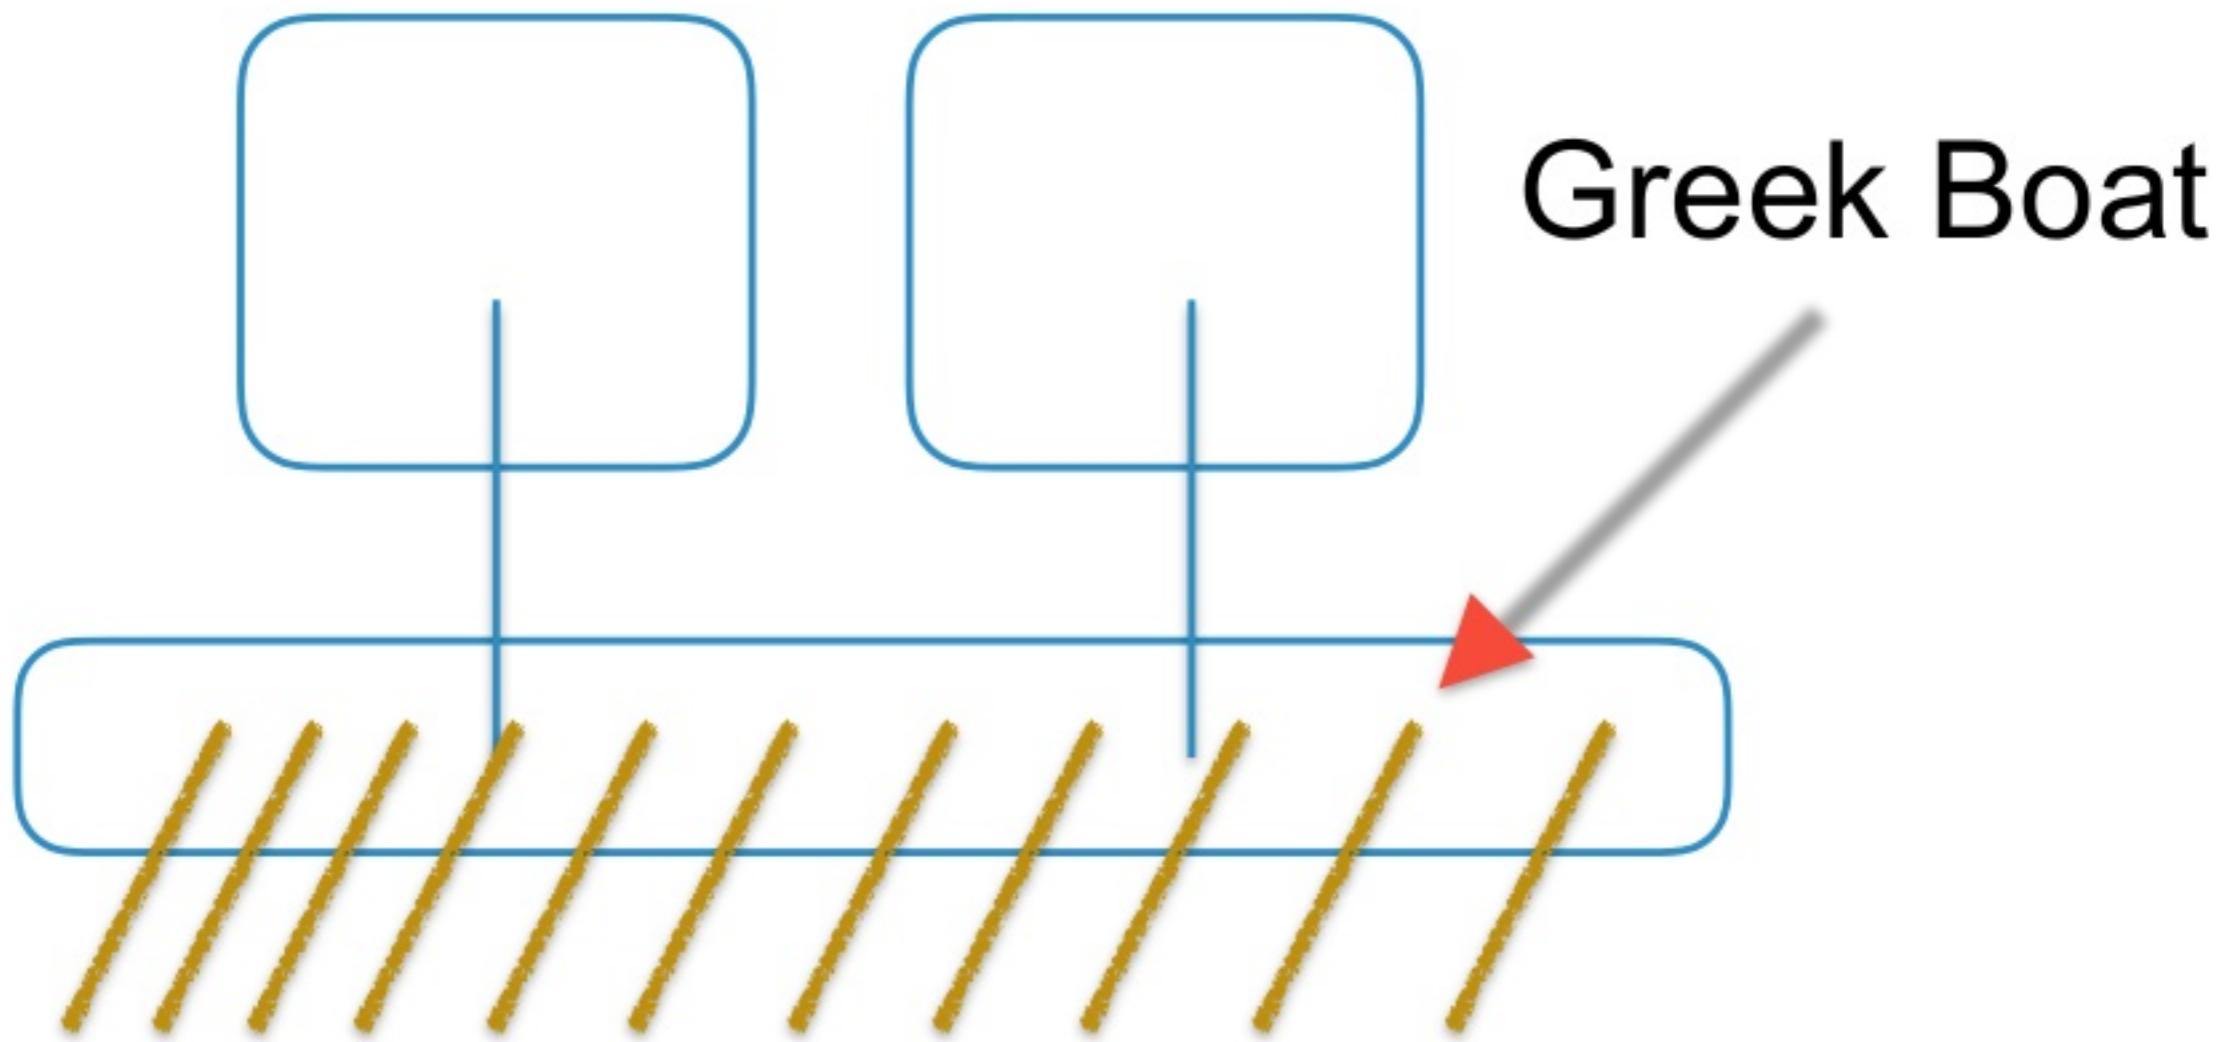


Have you heard of the
"Ship of Theseus?"

More Greek stuff ..



**When you replace all the parts of a thing Does it
Remain the Same?**



**When you replace all the parts of a thing Does it
Remain the Same?**



Hive Parser



Hive Optimization

SharkServer



Map-Reduce

APACHE
Spark™
Spark Execution



JDBC Results

**When you replace all the parts of a thing Does it
Remain the Same?**



Hive Parser



Hive Optimization

~~SharkServer~~
ThriftServer



Map-Reduce

APACHE
Spark™
Spark Execution



JDBC Results

**When you replace all the parts of a thing Does it
Remain the Same?**



Hive Parser

ThriftServer



Catalyst



Schema RDDs



Spark Execution



JDBC Results

**When you replace all the parts of a thing Does it
Remain the Same?**



Hive Parser



Catalyst

ThriftServer



DataFrames



Spark Execution



JDBC Results

**When you replace all the parts of a thing Does it
Remain the Same?**



Hive Parser

ThriftServer



Catalyst



Datasets



Spark Execution



JDBC Results

**When you replace all the parts of a thing Does it
Remain the Same?**

ThriftServer



Spark Parser



Catalyst



DataSets



Spark Execution



JDBC Results

Almost all Spark now

ThriftServer



JDBC Results

Connecting with Beeline (JDBC Client)

```
./bin/beeline  
dse beeline  
!connect jdbc:hive2://localhost:10000
```



Even More Hive!

Connect Tableau to Cassandra

Tableau - Book1

tab (ks.tab) (ks)

Connection: Live | Extract

Filters: 0 | Add...

Connected to Spark SQL

Server: 127.0.0.1

Schema: ks

Table: tab

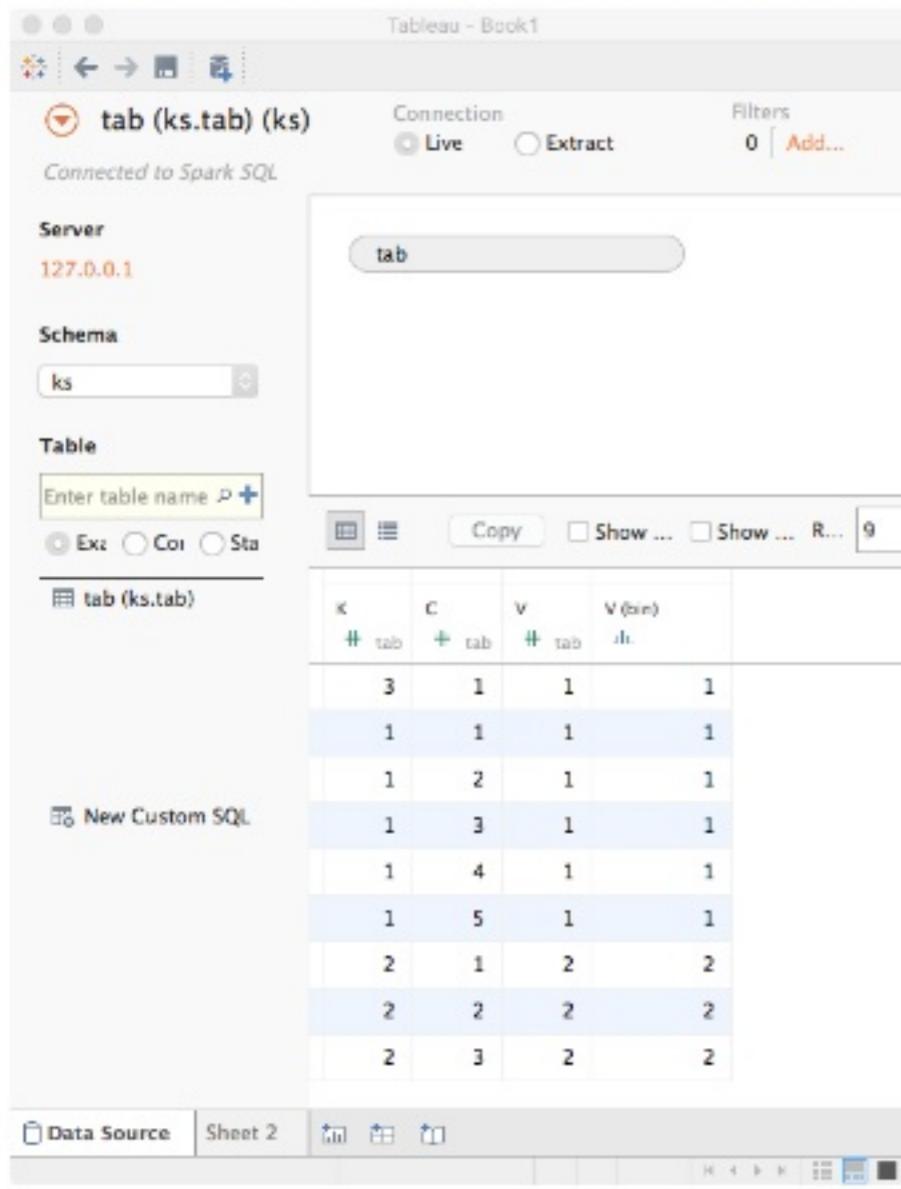
Enter table name: tab

Exz | Col | Sta

tab (ks.tab)

K	C	V	V (bin)
3	1	1	1
1	1	1	1
1	2	1	1
1	3	1	1
1	4	1	1
1	5	1	1
2	1	2	2
2	2	2	2
2	3	2	2

Data Source | Sheet 2



```
cqlsh:ks> SELECT * FROM tab ;
```



The Full JDBC/ODBC Ecosystem Can Connect to ThriftServer



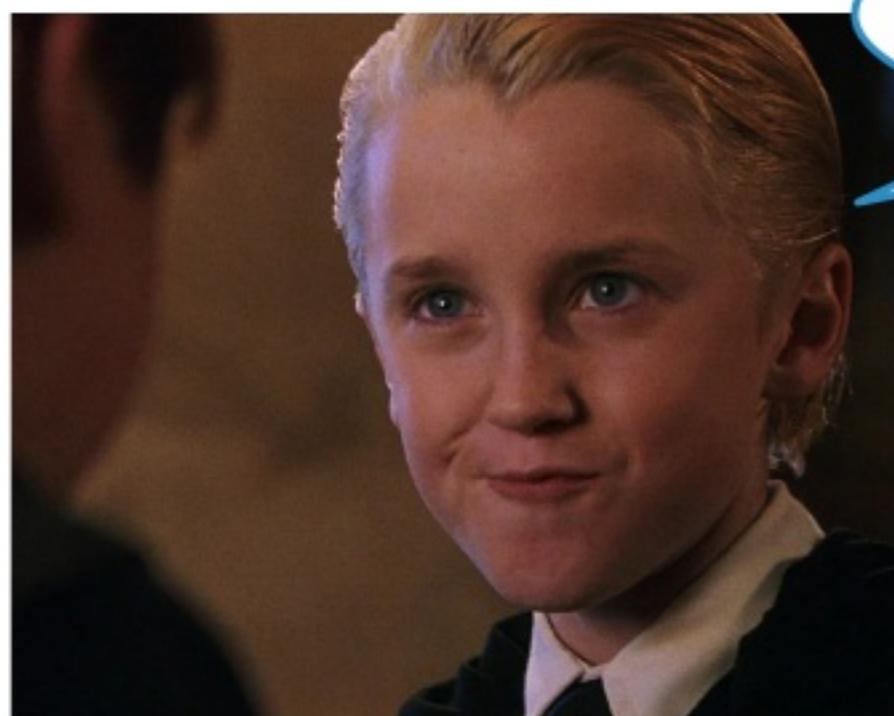
MicroStrategy



SAP BusinessObjects



Incremental Collect - Because BI Tools are Mean

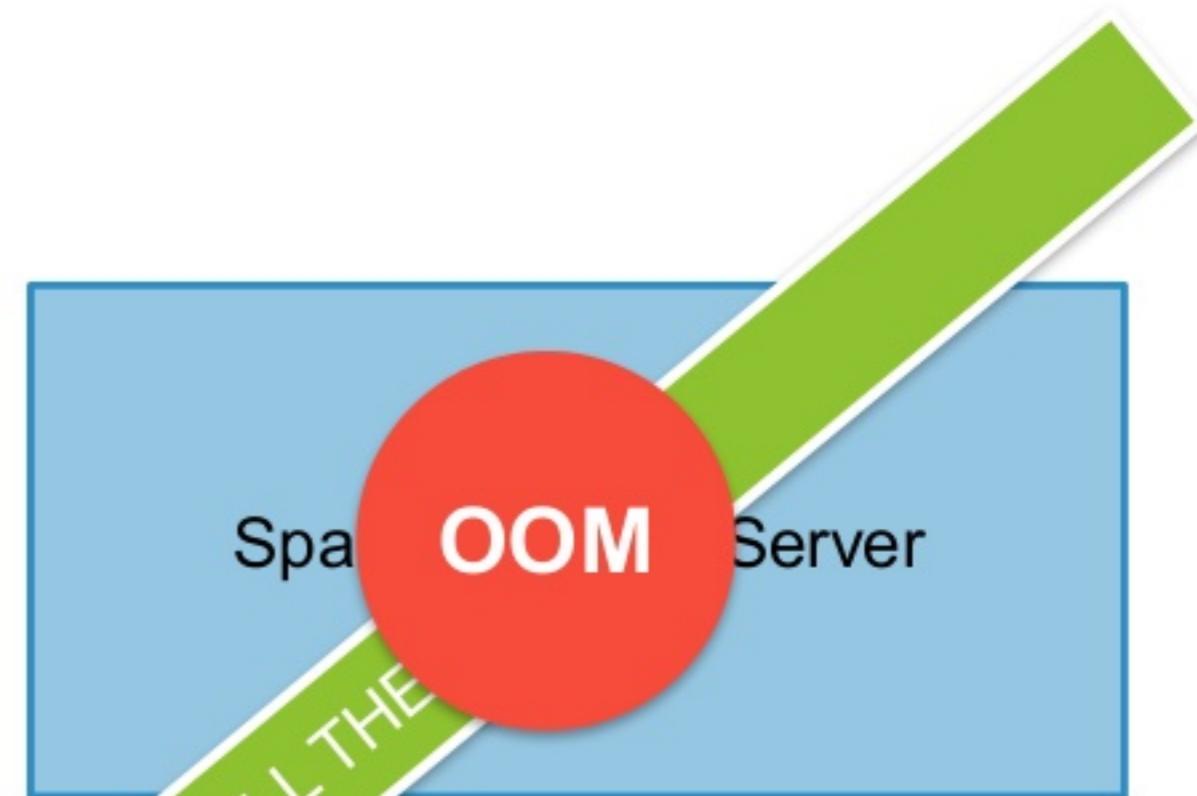


SELECT *
FROM TABLE

Spark Sql Thrift Server

ALL THE DATA

Incremental Collect - Because BI Tools are Mean



Incremental Collect - Because BI Tools are Mean



Spark Sql Thrift Server
`spark.sql.thriftServer.incrementalCollect=true`

Spark Partition 1

Spark Partition 2

Spark Partition 3

Incremental Collect - Because BI Tools are Mean



Spark Sql Thrift Server
`spark.sql.thriftServer.incrementalCollect=true`

Spark Partition 1

Spark Partition 2

Spark Partition 3

Incremental Collect - Because BI Tools are Mean



Spark Partition 1

Spark Sql Thrift Server
`spark.sql.thriftServer.incrementalCollect=true`

Spark Partition 2

ALL THE DATA

Spark Partition 3

Incremental Collect - Because BI Tools are Mean



Spark Partition 2

Spark Sql Thrift Server
`spark.sql.thriftServer.incrementalCollect=true`

Spark Partition 3

ALL THE DATA

Getting things done with SQL

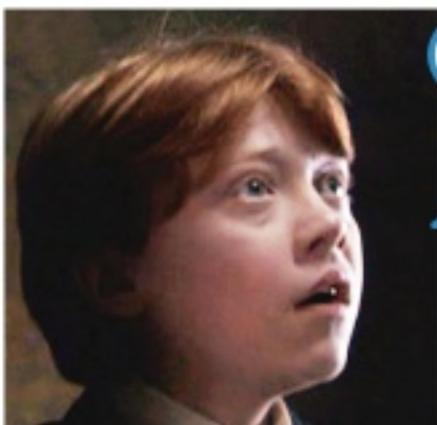
- Registering Sources
- Writing to Tables
- Examining Query Plans
- Debugging Predicate pushdowns
- Caching Views

Registering Sources using SQL

```
CREATE TEMPORARY VIEW words
    USING format.goes.here
    OPTIONS (
        key "value"
    )
```

Registering Sources using SQL

```
CREATE TEMPORARY VIEW words
  USING org.apache.spark.sql.cassandra
  OPTIONS (
    table "tab",
    keyspace "ks")
```



Not a single monad...

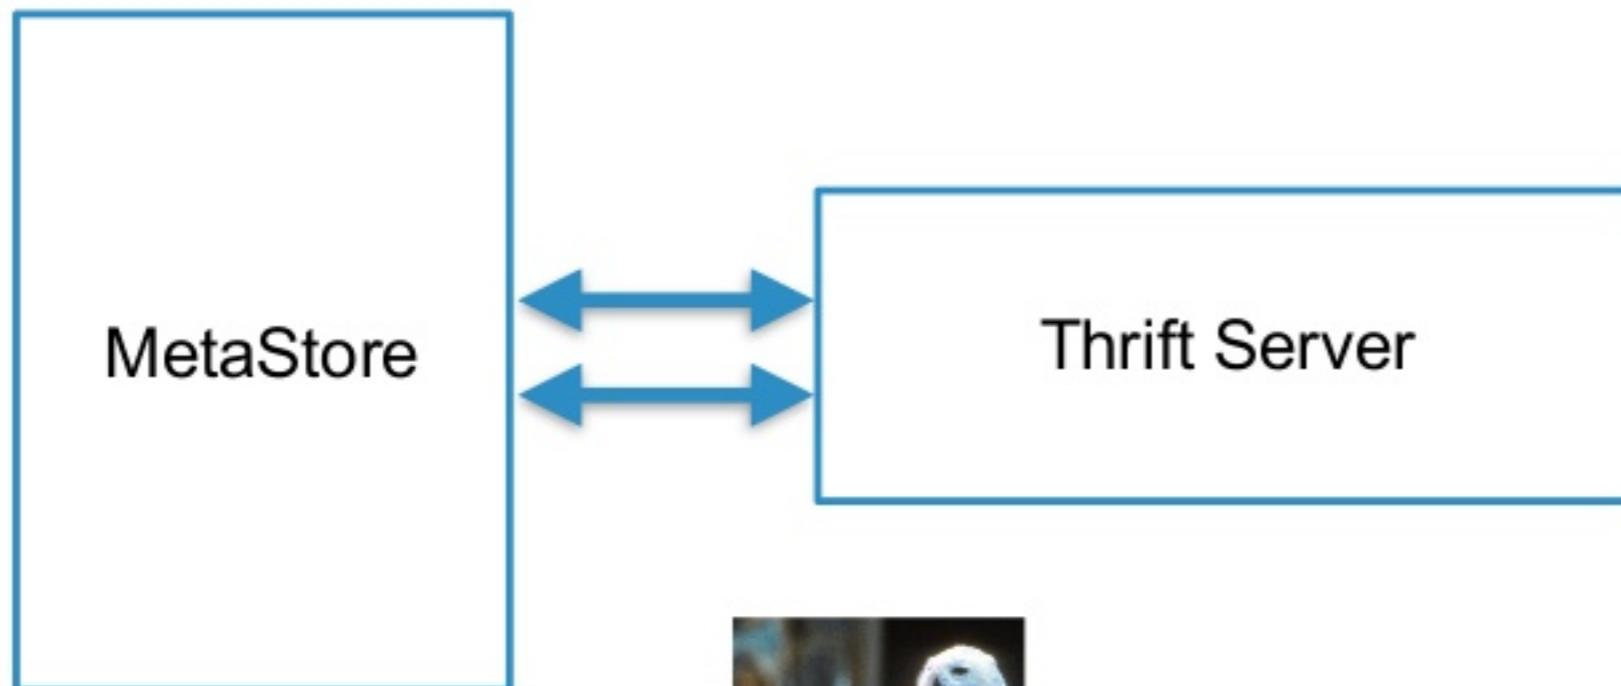
Registering Sources using SQL

```
CREATE TEMPORARY VIEW words
  USING org.apache.spark.sql.cassandra
  OPTIONS (
    table "tab",
    keyspace "ks")
```



CassandraSourceRelation

We Can Still Use a HiveMetaStore

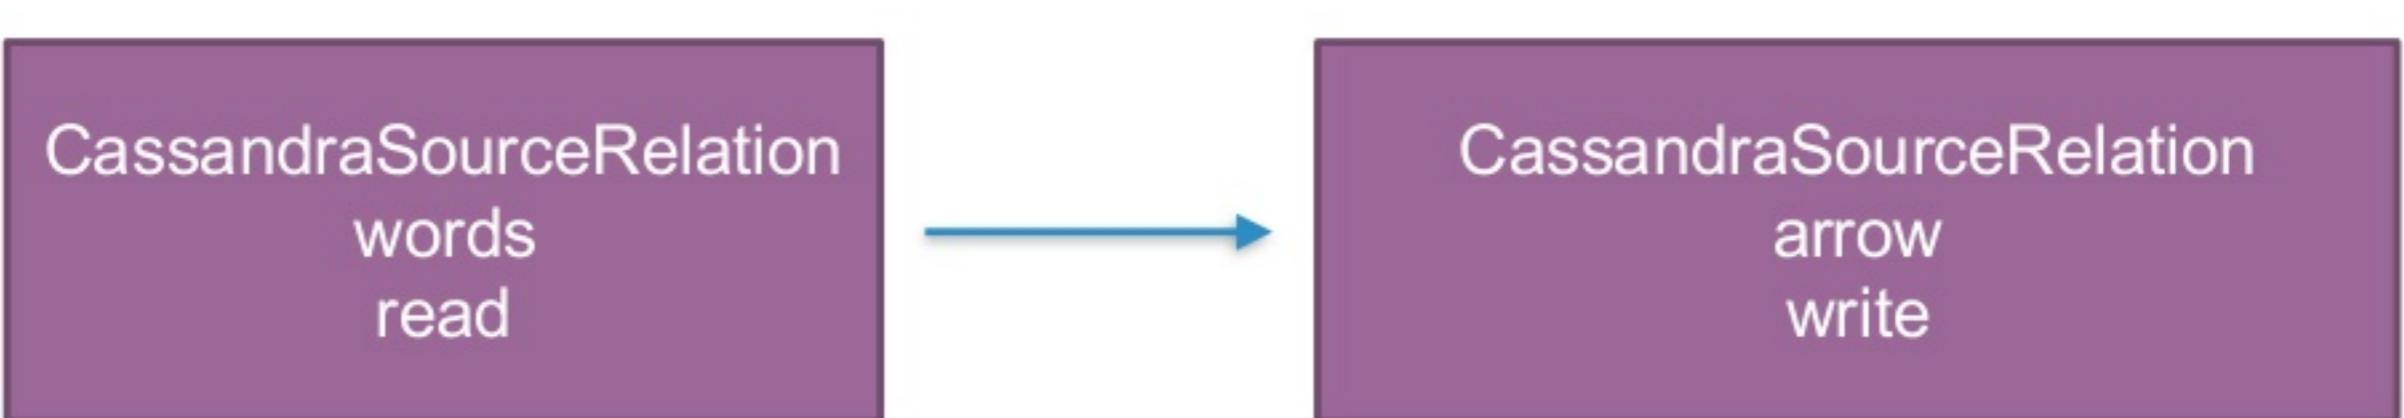


DATASTAX

DSE auto registers C* Tables in a C* based Metastore

Writing DataFrames using SQL

```
INSERT INTO arrow SELECT * FROM words;
```



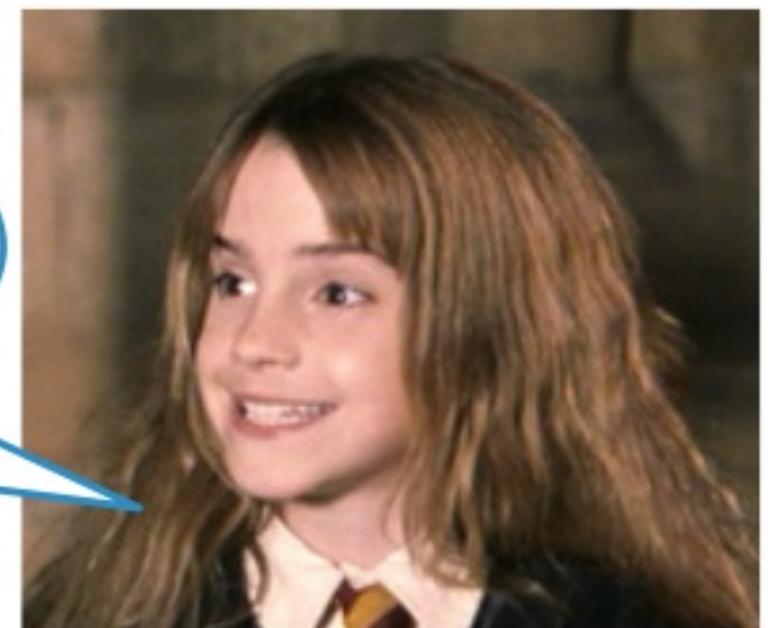
Explain to Analyze Query Plans

```
EXPLAIN SELECT * FROM arrow WHERE C > 2;
```

```
Scan org.apache.spark.sql.cassandra.CassandraSourceRelation@6069193a  
[k#18, c#19, v#20]
```

```
PushedFilters: [.IsNotNull(c), GreaterThan(c, 2)],  
ReadSchema: struct<k:int,c:int,v:int>
```

We can analyze the inside
of the Catalyst just
like with Scala/Java/...



Predicates get Pushed Down Automatically

```
EXPLAIN SELECT * FROM arrow WHERE C > 2;
```

```
Scan org.apache.spark.sql.cassandra.CassandraSourceRelation@6069193a  
[k#18, c#19, v#20]
```

```
PushedFilters: [.IsNotNull(c), GreaterThan(c,2)],  
ReadSchema: struct<k:int,c:int,v:int>
```



Predicates get Pushed Down Automatically

```
EXPLAIN SELECT * FROM arrow WHERE C > 2;
```

```
Scan org.apache.spark.sql.cassandra.CassandraSourceRelation@6069193a  
[k#18, c#19, v#20]
```

```
PushedFilters: [.IsNotNull(c), GreaterThan(c,2)],  
ReadSchema: struct<k:int,c:int,v:int>
```

CassandraSourceRelation

Filter [GreaterThan(c,2)]



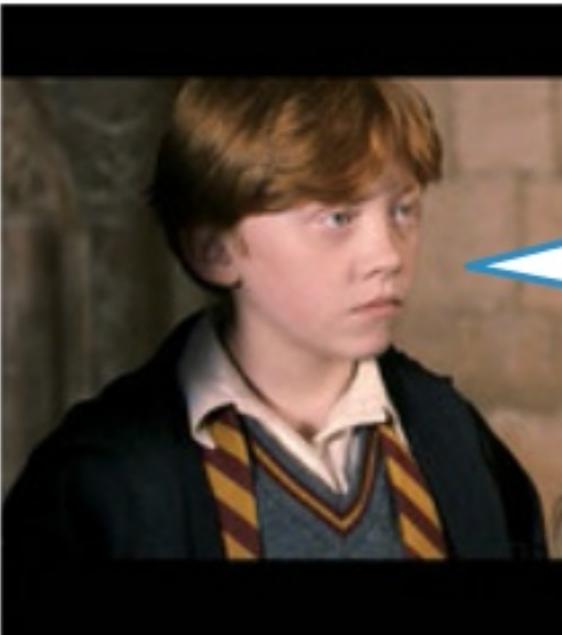
Internal Request to Cassandra: CQL
SELECT * FROM ks.bat WHERE C > 2

Automatic
Pushdowns!



Common Cases where Predicates Don't Push

```
SELECT * from troubles WHERE c < '2017-05-27'  
*Filter (cast(c#76 as string) < 2017-05-27)  
+- *Scan CassandraSourceRelation@53e82b30 [k#75,c#76,v#77]  
    PushedFilters: [IsNotNull(c)],  
    ReadSchema: struct<k:int,c:date,v:int>
```



Why is my date
clustering column not being
pushed down.

Common Cases where Predicates Don't Push

```
SELECT * from troubles WHERE c < '2017-05-27'  
*Filter (cast(c#76 as string) < 2017-05-27)  
+- *Scan CassandraSourceRelation@53e82b30 [k#75,c#76,v#77]  
    PushedFilters: [IsNotNull(c)],  
    ReadSchema: struct<k:int,c:date,v:int>
```

CassandraSourceRelation

Filter [LessThan(c,'2017-05-27')]



Common Cases where Predicates Don't Push

```
SELECT * from troubles WHERE c < '2017-05-27'  
*Filter (cast(c#76 as string) < 2017-05-27)  
+- *Scan CassandraSourceRelation@53e82b30 [k#75,c#76,v#77]  
    PushedFilters: [IsNotNull(c)],  
    ReadSchema: struct<k:int,c:date,v:int>
```

CassandraSourceRelation
ReadSchema:
struct<k:int,c:date,v:int>



Filter [LessThan(c,'2017-05-27')]
Date != String

Make Sure we Cast Correctly

```
EXPLAIN SELECT * from troubles WHERE  
c < cast('2017-05-27' as date);
```

*Scan C*Relation PushedFilters: [IsNotNull(c), LessThan(c, 2017-05-27)]

CassandraSourceRelation
ReadSchema:
struct<k:int,c:date,v:int>



Filter [LessThan(c,Date('2017-05-27'))]
Date == Date

Make Sure we Cast Correctly

```
EXPLAIN SELECT * from troubles WHERE  
c < cast('2017-05-27' as date);
```

*Scan C*Relation PushedFilters: [IsNotNull(c), LessThan(c, 2017-05-27)]

CassandraSourceRelation
ReadSchema:
struct<k:int,c:date,v:int>

Filter [LessThan(c,Date('2017-05-27'))]

Automatic
Pushdowns!



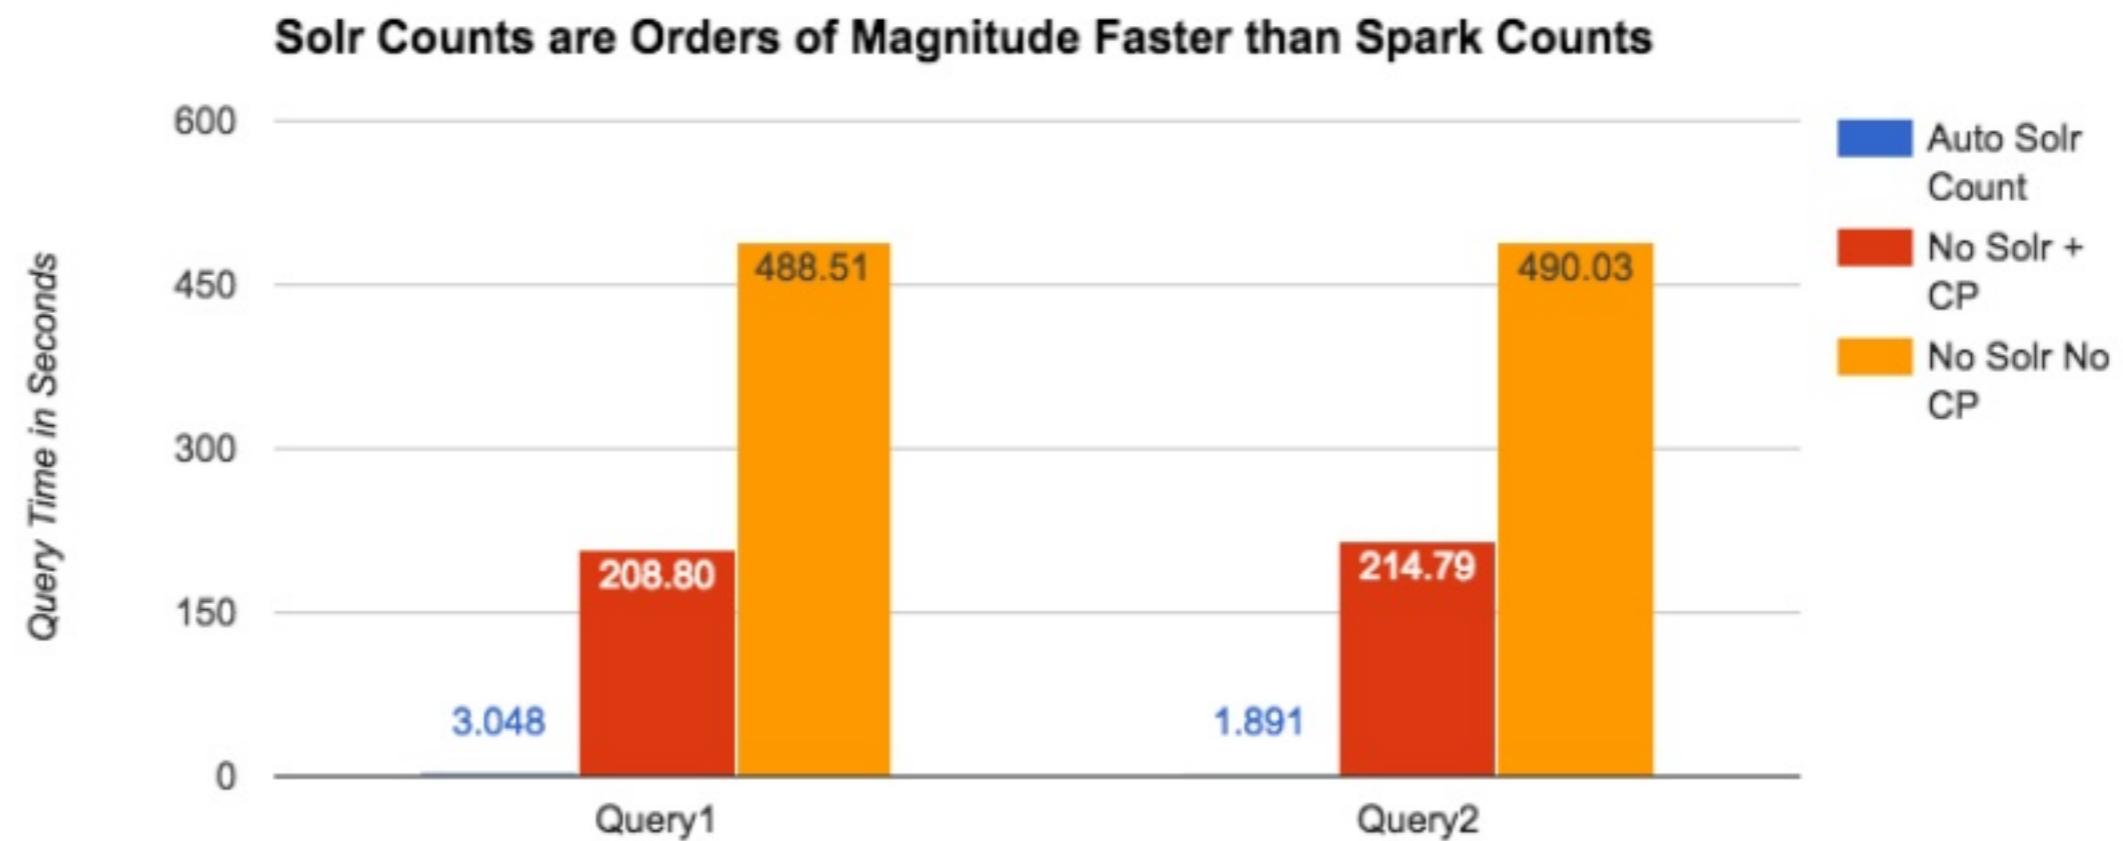
DSE Search Automatic Pushdowns!

```
EXPLAIN SELECT * from troubles WHERE v < 6;
```

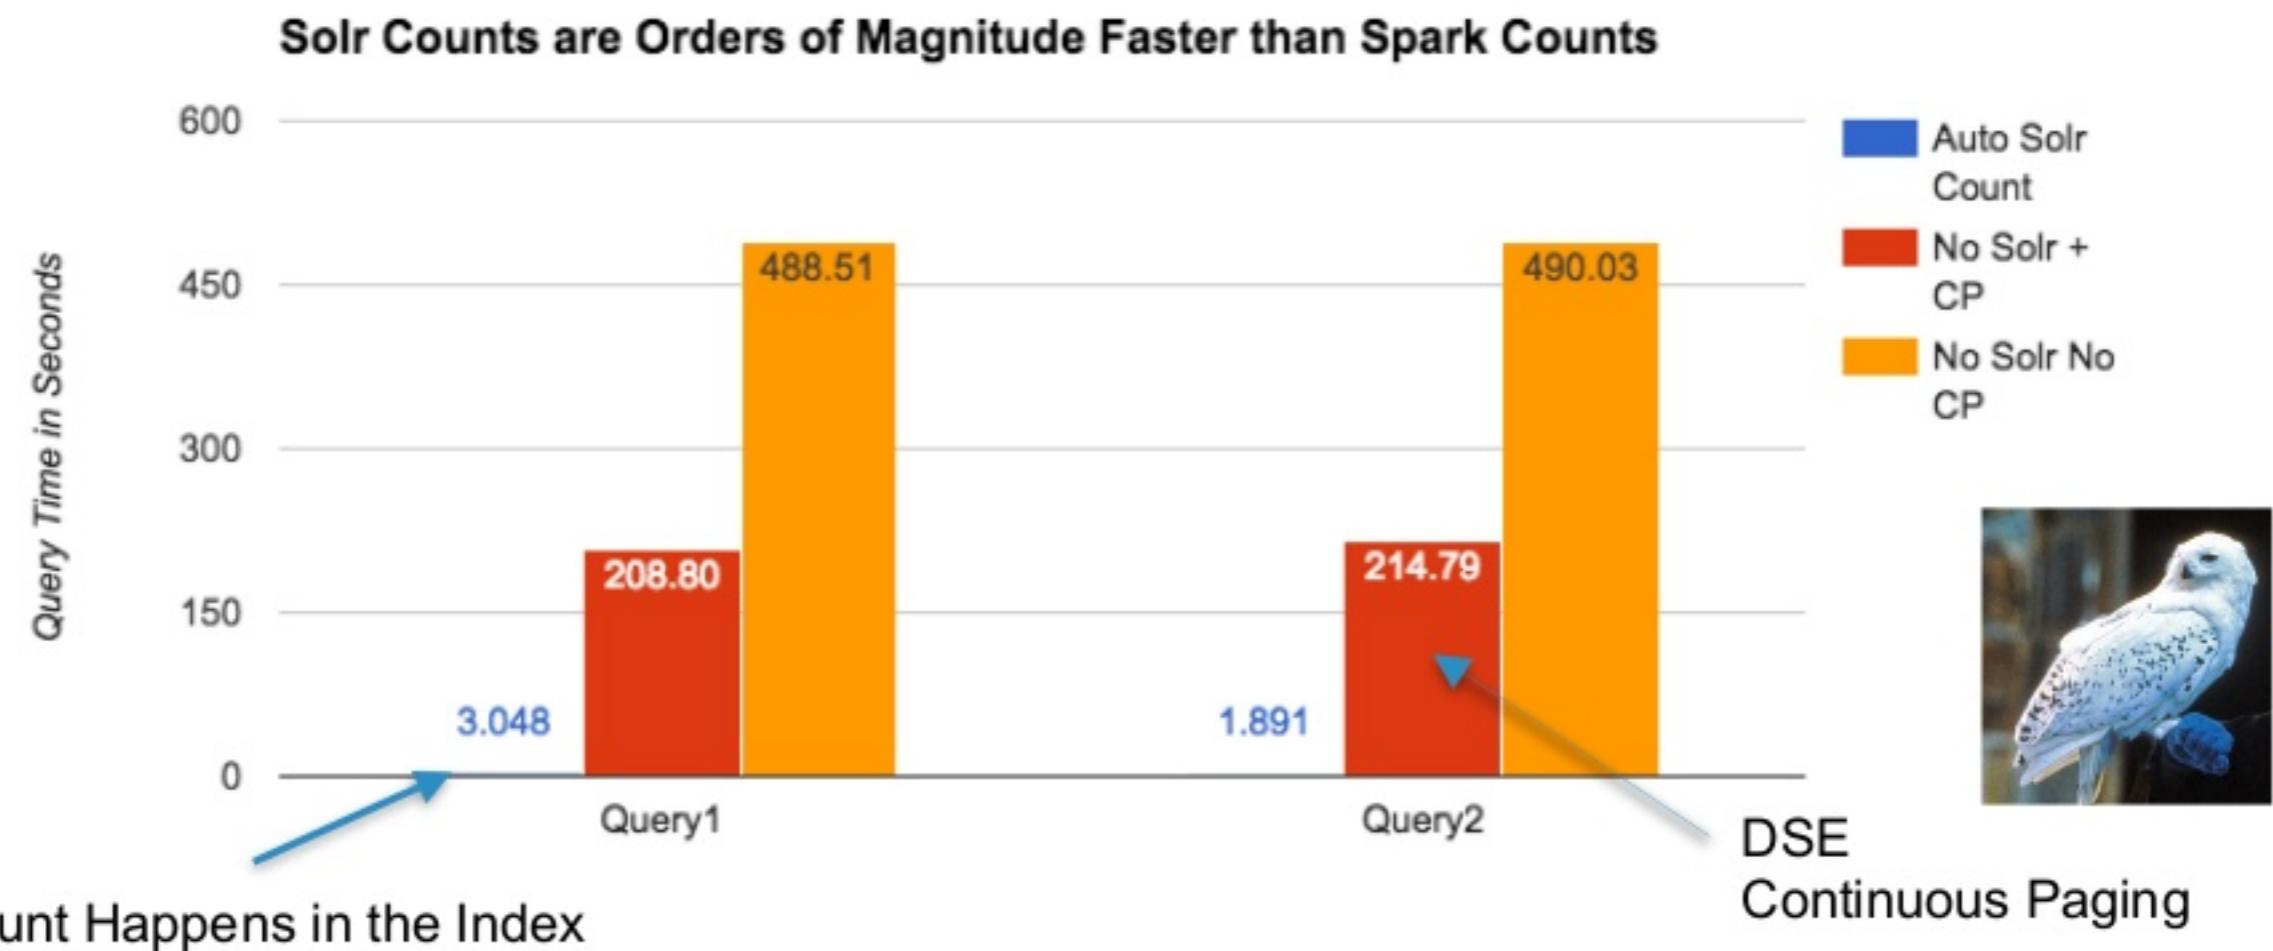
*Scan C*Relation PushedFilters: [IsNotNull(v), LessThan(v, 6)]



DSE Search Automatic Pushdowns!

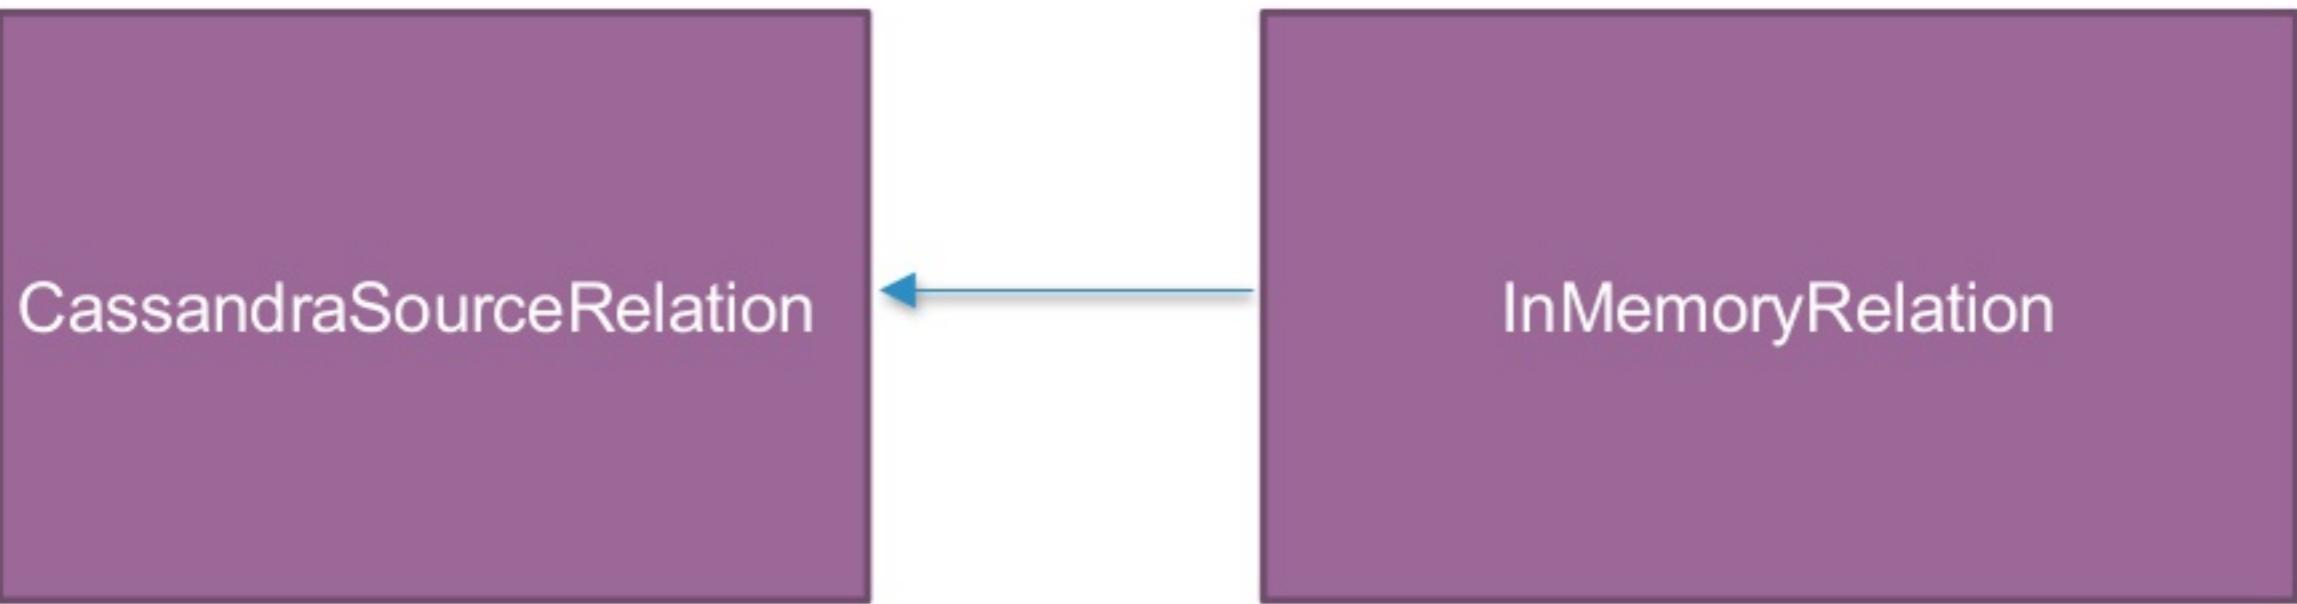


DSE Search Automatic Pushdowns!



Cache a whole table

```
CACHE TABLE ks.tab;  
explain SELECT * FROM ks.tab;  
== Physical Plan ==  
InMemoryTableScan [k#0, c#1, v#2]  
:  +- InMemoryRelation StorageLevel(disk, memory, deserialized, 1 replicas), `ks`.`tab`  
:    :  +- *Scan CassandraSourceRelation
```



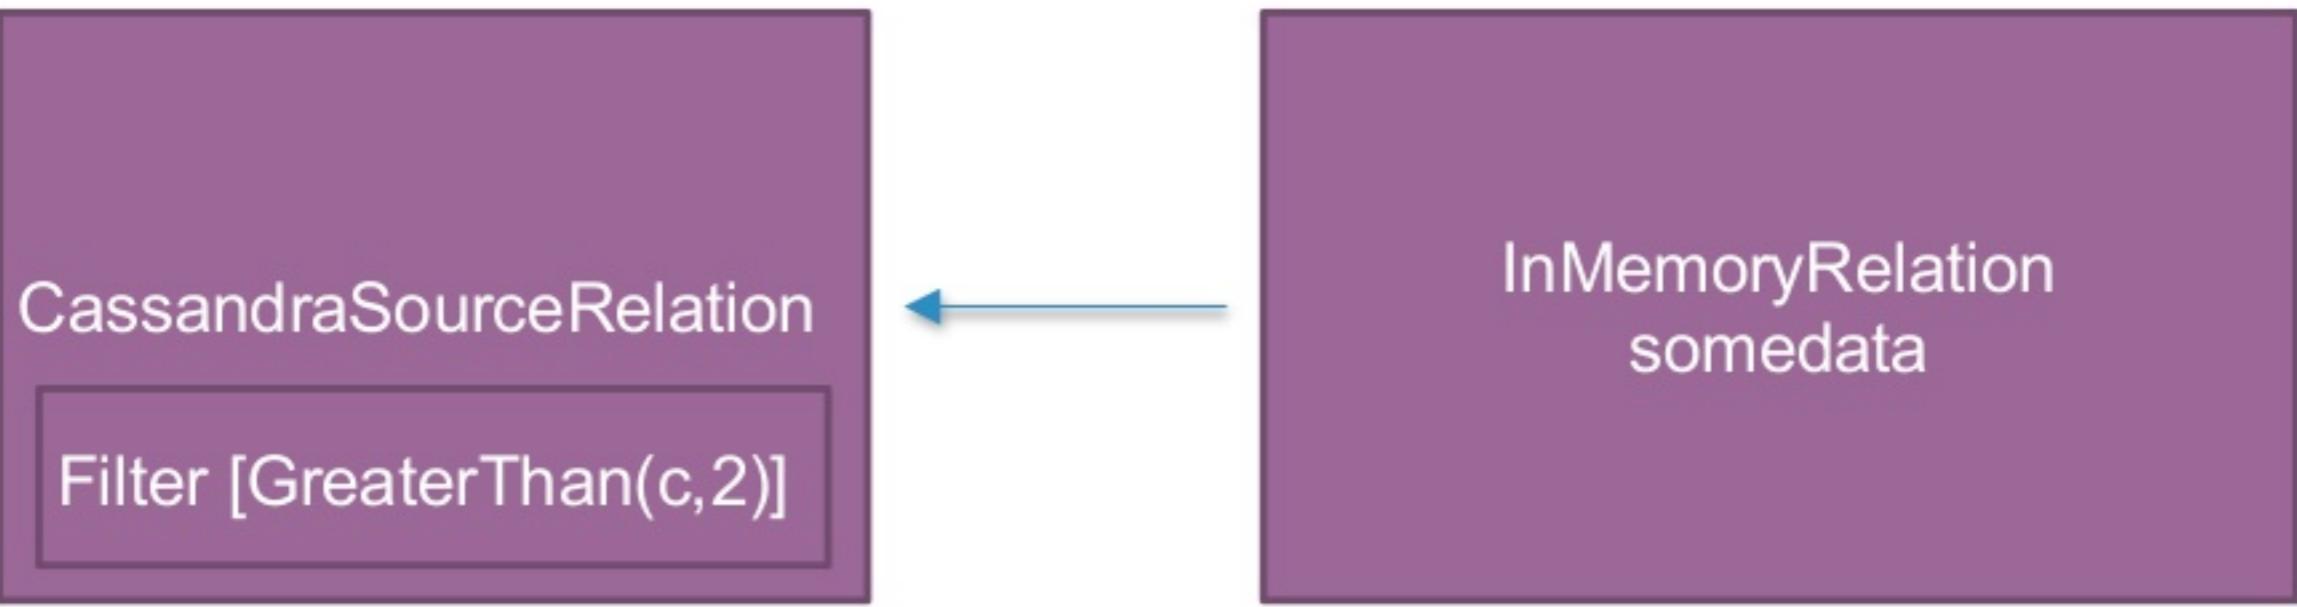
Uncache

```
UNCACHE TABLE ks.tab;  
explain SELECT * FROM ks.tab;  
== Physical Plan ==  
*Scan CassandraSourceRelation
```

CassandraSourceRelation

Cache a fraction of Data

```
CACHE TABLE somedata SELECT * FROM ks.tab WHERE c > 2;  
explain SELECT * from somedata;  
== Physical Plan ==  
InMemoryTableScan  
: +- InMemoryRelation `somedata`  
:   : +- *Scan CassandraSourceRelation PushedFilters: [IsNotNull(c), GreaterThan(c,2)]
```



Let this be a starting point

- <https://github.com/datastax/spark-cassandra-connector>
- https://github.com/datastax/spark-cassandra-connector/blob/master/doc/14_data_frames.md
- <https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-sql-thrift-server.html>
- http://docs.datastax.com/en/dse/5.1/dse-dev/datastax_enterprise/spark/sparkSqlThriftServer.html
- <https://spark.apache.org/docs/latest/sql-programming-guide.html#distributed-sql-engine>
- <https://www.datastax.com/dev/blog/dse-5-1-automatic-optimization-of-spark-sql-queries-using-dse-search>
- <https://www.datastax.com/dev/blog/dse-continuous-paging-tuning-and-support-guide>





Thank You.

<http://www.russellsipitzer.com/>
[@RussSpitzer](https://twitter.com/RussSpitzer)

Come chat with us at DataStax Academy:
<https://academy.datastax.com/slack>